

Space-Efficient Data Structures

Francisco Claude
Gonzalo Navarro



continues...


The Goal

Design data structures that:

- Have a small memory footprint
- Support fast queries


Web Graphs

- UK-Union-2006-06-2007-05
- Nodes: 133,633,040
- Edges: 5,507,679,822
- Plain representation: 22GB



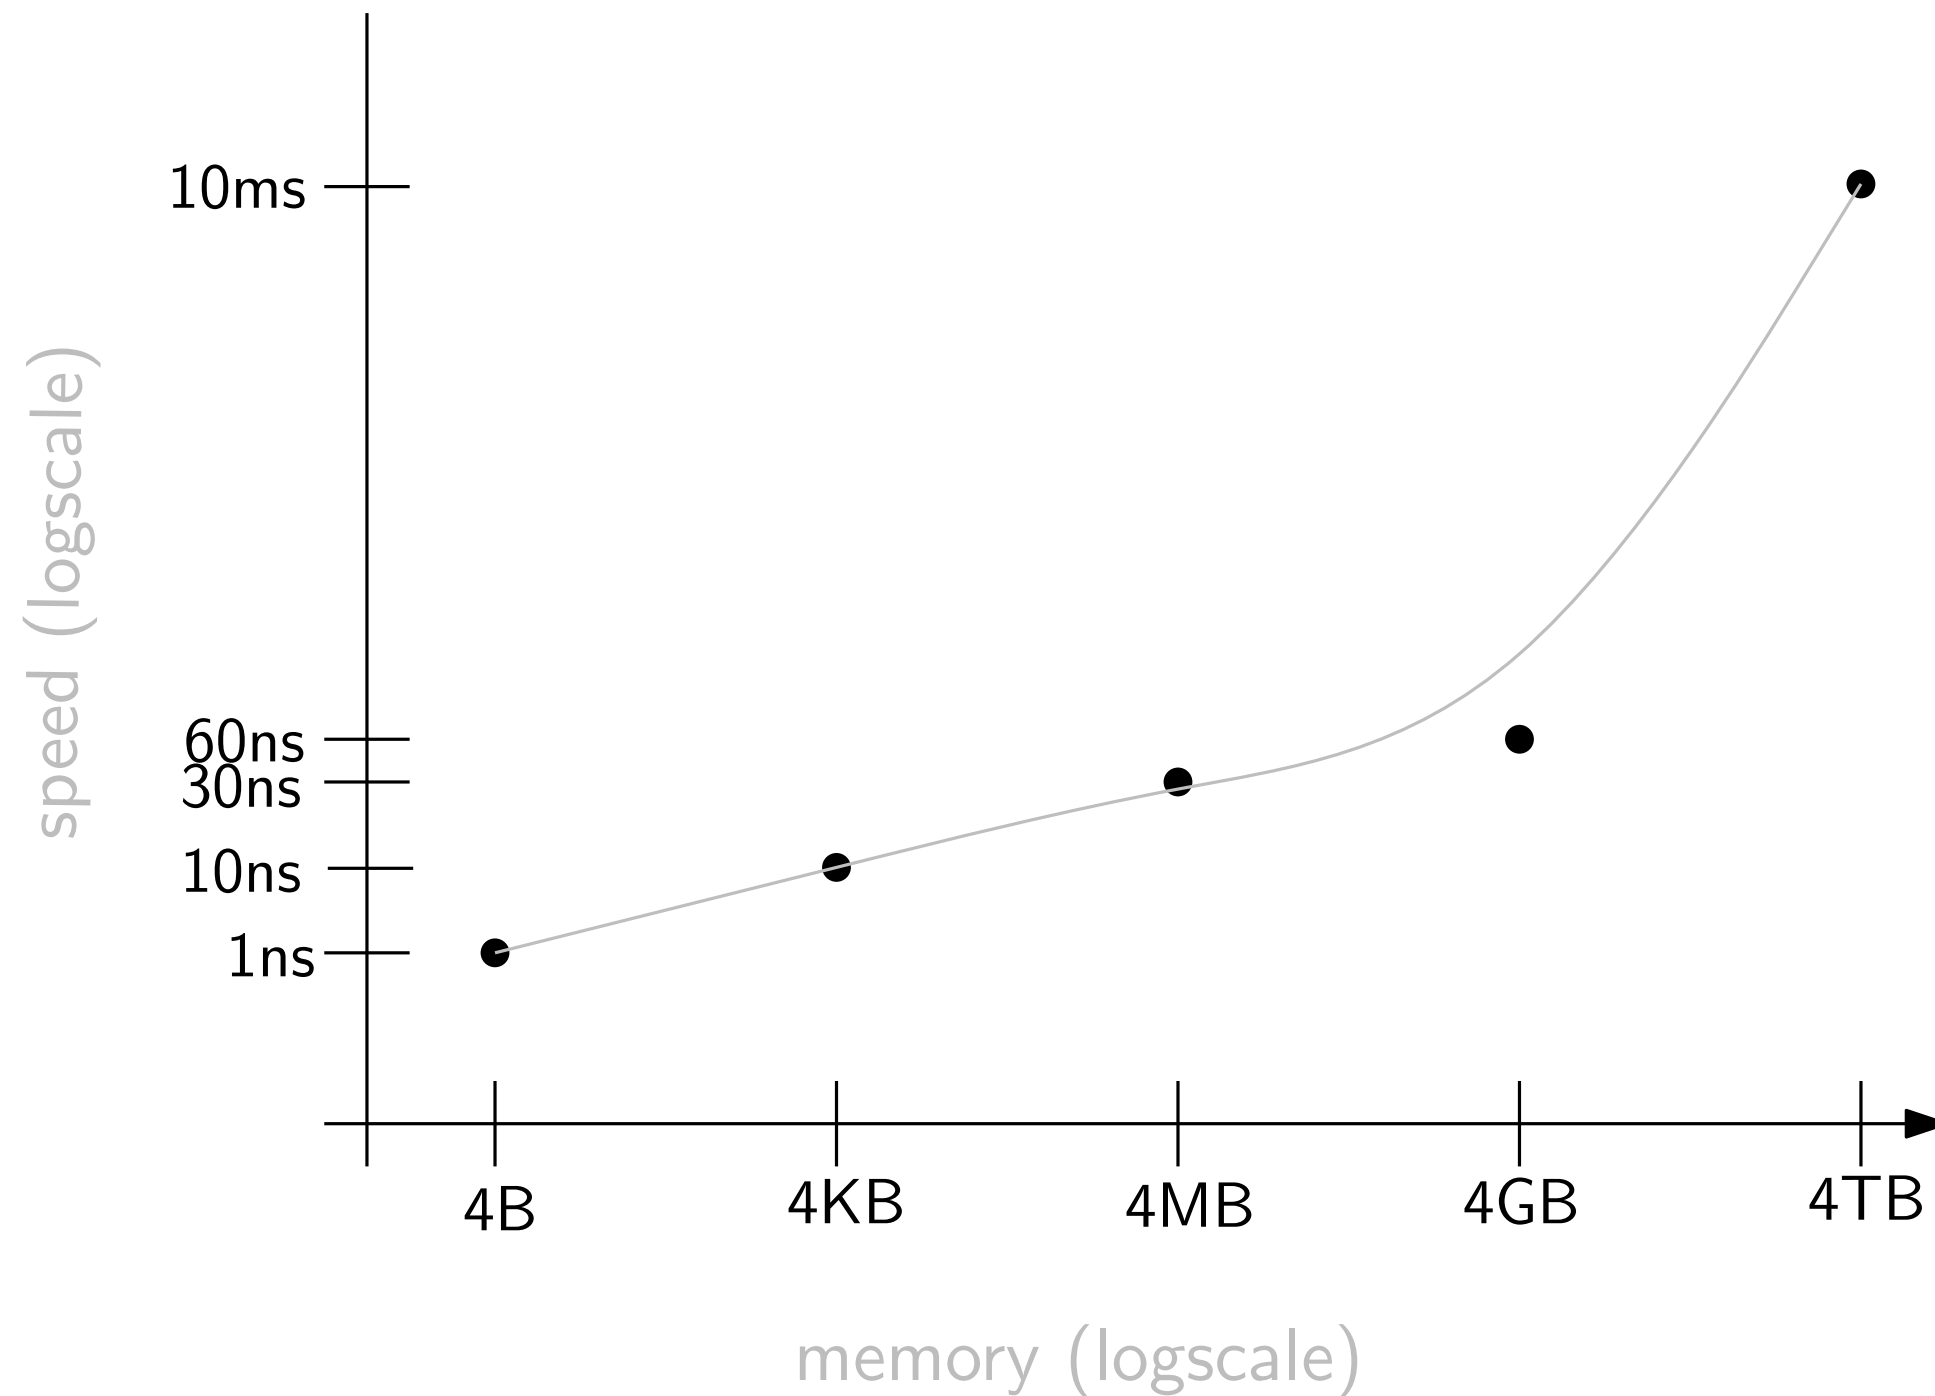
Web Graphs

- UK-Union-2006-06-2007-05
- Nodes: 133,633,040
- Edges: 5,507,679,822
- Plain representation: 22GB

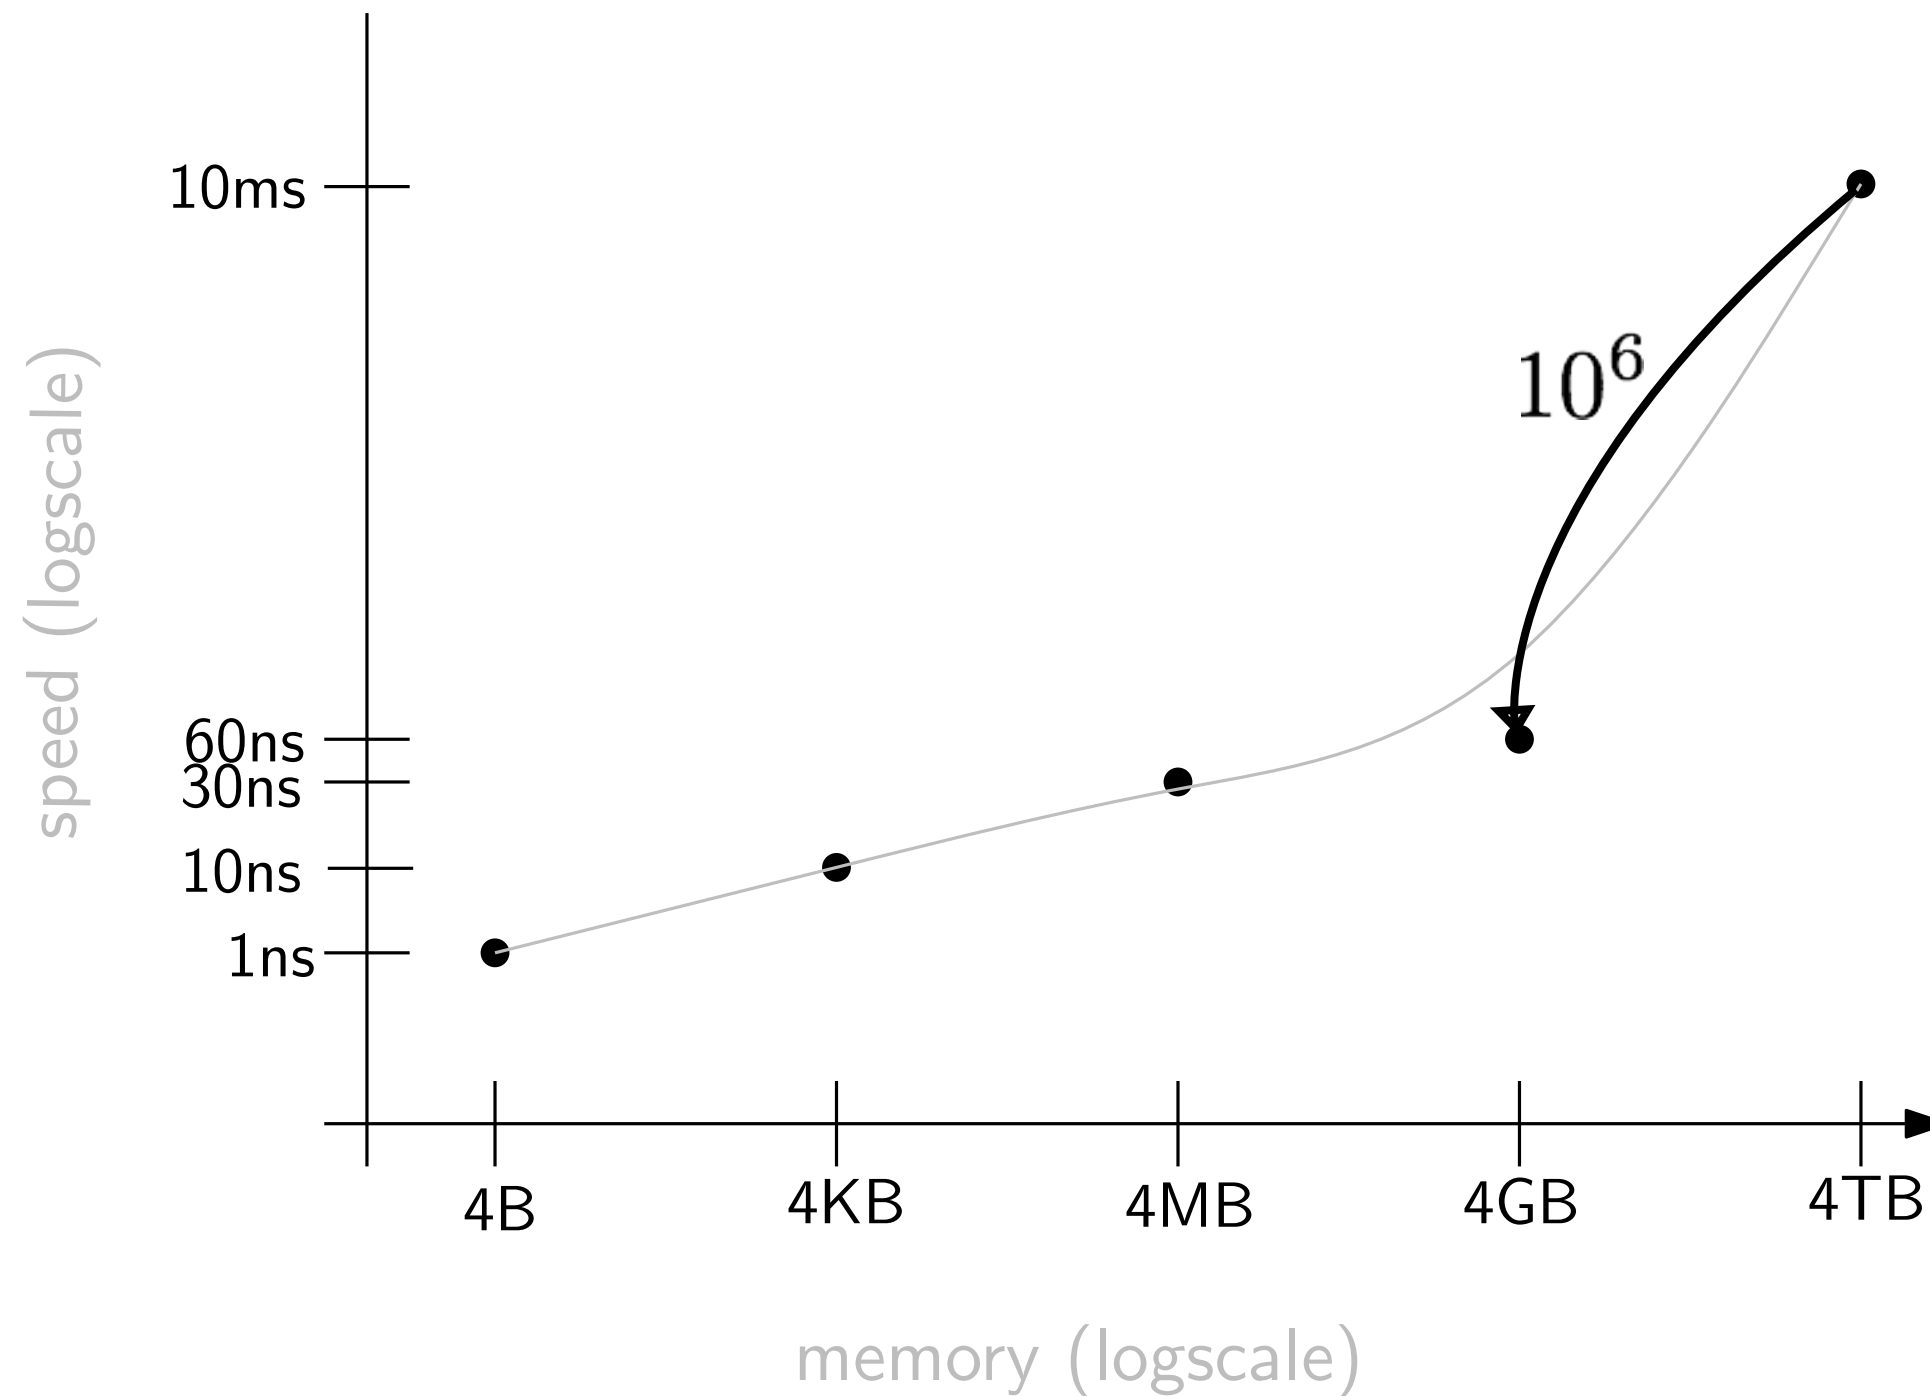


We can get it
down to <2GB!!

Motivation



Motivation



Outline

- Representing information
- Bitmaps
 - Trees (intro)
- Sequences (and permutations)
- Applications
- Conclusions

The model

- Word-RAM model
 - RAM of size n
 - We can manipulate $w = \Theta(\log n)$ bits at the time
 - CPU with $O(1)$ registers
 - Operations $+$, $-$, $*$, $/$, $<<$, $>>$ take constant time -- we can address with the result

Arrays

- Store elements addressed by an index
- Support efficient access
- Ideally, support some sort of mutation

What we are used to

```
uint *a = (uint*)malloc(sizeof(uint) * n);  
... a[i] ...
```

```
uint *a = new uint[n];  
... a[i] ...
```

```
a := make([]uint32, n)  
... a[i] ...
```

...

What we are used to

```
uint *a = (uint*)malloc(sizeof(uint) * n);  
... a[i] ...
```

```
uint *a = new uint[n];  
... a[i] ...
```

```
a := make([]uint32, n)  
... a[i] ...
```

... We use 32/64 bits per element

What if all values are small?

- We may not need 32/64 bits per element
- Say the maximum value is m
- We can use $\lceil \log_2(m + 1) \rceil$ bits per element

What if all values are small?

- We may not need 32/64 bits per element
- Say the maximum value is m
- We can use $\lceil \log_2(m + 1) \rceil$ bits per element

```
inline void SetField(cds_word *A, const cds_word len, const cds_word index, const cds_word x)
{
    if (len == 0) {
        return;
    }
    cds_word i = index * len / kWordSize, j = index * len - i * kWordSize;
    cds_word mask = ((j + len) < kWordSize ? ~(cds_word)0 << (j + len) : 0)
                    | ((kWordSize - j) < kWordSize ? ~(cds_word)0 >> (kWordSize - j) : 0);
    A[i] = (A[i] & mask) | x << j;
    if (j + len > kWordSize) {
        mask = ((~(cds_word)0) << (len + j - kWordSize));
        A[i + 1] = (A[i + 1] & mask) | x >> (kWordSize - j);
    }
}
```

Space Required (2^{32} elements)

Number of bits per element	Total space (GBs)
2	1
4	2
8	4
16	8
32	16
64	32

Arrays in LIBCDS

```
cds_word *values = new cds_word[n];  
for (cds_word i = 0; i < n; i++) {  
    values[i] = ComputeValueAt(i);  
}  
Array *A = Array::Create(values, n);  
A->SetField(0, 1);  
assert (A->GetField(0) == 1);
```

```
Array *A = Array::Create(n, bits);  
for (cds_word i = 0; i < n; i++) {  
    A->SetField(i, ComputeValueAt(i));  
}
```


Compression

- What happens if some values tend to repeat a lot?
- Can we do better?
- Yes, we can assign shorter codes to the most frequent elements (Huffman for example)

$$H_0(S) = \sum_{c \in \Sigma} \frac{n_c}{n} \log_2 \frac{n}{n_c}$$

Compression

- Sequence $S = \text{aaabbcaabbcaaad}$

$$H_0(S) = \sum_{c \in \Sigma} \frac{n_c}{n} \log_2 \frac{n}{n_c}$$

$$H_0(S) = 1.5919$$

symb	freq
a	9
b	4
c	2
d	1


Compression

- Sequence $S = \text{aaabbcaabbcaaad}$

$$H_0(S) = \sum_{c \in \Sigma} \frac{n_c}{n} \log_2 \frac{n}{n_c}$$

$$H_0(S) = 1.5919$$

symb	freq
a	9
b	4
c	2
d	1




Compression

- Sequence $S = \text{aaabbcaabbcaaad}$

$$H_0(S) = \sum_{c \in \Sigma} \frac{n_c}{n} \log_2 \frac{n}{n_c}$$

$$H_0(S) = 1.5919$$

symb	freq
a	9
b	4
c	2
d	1



$$H_k(S) = \sum_{|T|=k} |S^T| H_0(S^T)$$

Huffman


- Huffman's algorithm
 - Every element is a tree
 - Iteratively, take the two least frequent trees and merge them
 - Stop when there is only one tree

Huffman

symb	freq
a	9
b	4
c	2
d	1



Huffman




symb	freq
a	9
b	4
c	2
d	1

Huffman




symb	freq
a	9
b	4
c	2
d	1

Huffman



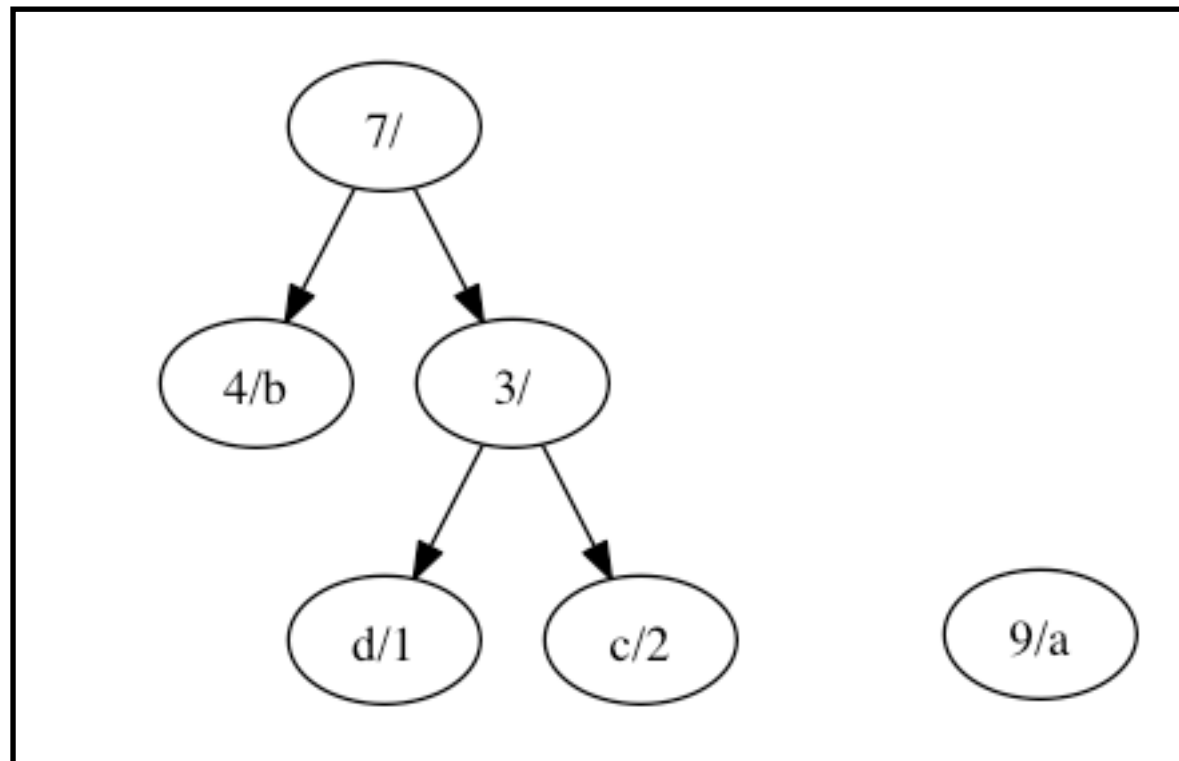
symb	freq
a	9
b	4
c	2
d	1

Huffman




symb	freq
a	9
b	4
c	2
d	1

Huffman




symb	freq
a	9
b	4
c	2
d	1

Huffman



symb	freq
a	9
b	4
c	2
d	1

Huffman



symb	freq
a	9
b	4
c	2
d	1

symb	code
a	1
b	00
c	011
d	010

Huffman

- Sequence $S = \text{aaabbcaaabbcaaad}$

| | | 00000 | | ...

Huffman

- Sequence $S = \text{aaabbcaaabbcaaad}$

1 1 1 0 0 0 0 0 1 1 ...

sy mb	code
a	1
b	00
c	011
d	010


Huffman & Random access

- We can't access an arbitrary position
- One simple solution is to sample the starting position every k elements
- This allows to access in $O(k)$

Huffman + bitmap

- To access position i , we just do


$\text{decode}(H[\text{select}(l, i) \dots \text{select}(l, i+1)-1])$



Huffman + bitmap

- To access position i , we just do

$\text{decode}(H[\text{select}(l, i) \dots \text{select}(l, i+1)-1])$



Huffman & Random access

- We could mark the beginning of each code with a 1 in a separate bitmap that runs in parallel
- If we could find the i -th 1 in the bitmap in constant time, we would be able to access the i -th code.
- This motivates the following...

Huffman & Random access

- Another option is DACs
- In the first level, write down the first bit of each code
- In a bitmap in parallel, mark which codes continue to the next level
- Continue recursively with the next levels

DACs

- Sequence $S = \text{aaabbcaaabbcaaad}$

| | | 0 0 0 0 0 | | ...

1	1	1	0	0	0	...		
0	0	0	1	1	1			
0	0	1	...					
0	0	1						
1	...							
0								

DACs

- Sequence $S = \text{aaabbcaaabbcaaad}$

| | | 0 0 0 0 0 | | ...

1	1	1	0	0	0	...		
0	0	0	1	1	1			
0	0	1	...					
1	...							
0								

sympb	code
a	1
b	00
c	011
d	010

Bitmaps

- $\text{access}(i)$: retrieve i -th bit
- $\text{rank}(0/1, i)$: count how many 0/1s appear up to position i
- $\text{select}(0/1, j)$: find the j -th occurrence of 0/1

$$B = \boxed{1001110000111}100000\boxed{0}$$

$\text{rank}(13) = 7$

$\text{access}(19) = 0$

$\text{select}(8) = 14$

Small bitmaps

- We will build the solution bottom-up
- Consider bitmaps of size $O(\log n)$

Small bitmaps

- Access is trivial using \ll and \gg ($v \& (1 \ll i)$)
- Rank: store all possible answers for bitmaps of length $\frac{\log n}{2}$!

Small bitmaps

- Access is trivial using \ll and \gg ($v \& (1 \ll i)$)
- Rank: store all possible answers for bitmaps of length $\frac{\log n}{2}$!

$2^{\frac{\log n}{2}}$ bitmaps of length $\frac{\log n}{2}$

Small bitmaps

- Access is trivial using \ll and \gg ($v \& (1 \ll i)$)
- Rank: store all possible answers for bitmaps of length $\frac{\log n}{2}$!

$2^{\frac{\log n}{2}}$ bitmaps of length $\frac{\log n}{2}$
 $\frac{\log n}{2}$ queries for each, and the
answer takes $\log \log n$ bits

Small bitmaps

- Access is trivial using \ll and \gg ($v \& (1 \ll i)$)
- Rank: store all possible answers for bitmaps of length $\frac{\log n}{2}$!

$2^{\frac{\log n}{2}}$ bitmaps of length $\frac{\log n}{2}$
 $\frac{\log n}{2}$ queries for each, and the
 answer takes $\log \log n$ bits

Total Space: $\frac{\sqrt{n} \log n \log \log n}{2}$

Small bitmaps

001110011
rank(1,8)

B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3

Small bitmaps

001110011
rank(1,8)

B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3

Small bitmaps

001110011
rank(1,8)

B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3

1

Small bitmaps

001110011
rank(1,8)

1

B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3

Small bitmaps

001110011
rank(1,8)

1

B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3

Small bitmaps

001110011
rank(1,8)

1

B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3

2

Small bitmaps

001110011
rank(1,8)

1+2

B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3

2

Small bitmaps

001110011
rank(1,8)

1+2


B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3

Small bitmaps

001110011
rank(1,8)

1+2

B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3




Small bitmaps

001110011
rank(1,8)

1 + 2 + 1

B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3



Small bitmaps

001110011
rank(1,8)

1+2+1

B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3

Small bitmaps

001110011
rank(1,8)

$$1 + 2 + 1 = 4$$

B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3

Small bitmaps

001110011
rank(1,8)

$$1 + 2 + 1 = 4$$

B	p=1	p=2	p=3
000	0	0	0
001	0	0	1
010	0	1	1
011	0	1	2
100	1	1	1
101	1	1	2
110	1	2	2
111	1	2	3

For a word of size $c \log n$ we do $2c$ lookups

Small bitmaps

- How big is the table?

$\frac{\log n}{2}$	KBs
8	0.75
16	512

Small bitmaps

- How big is the table?

Total Space: $\frac{\sqrt{n} \log n \log \log n}{2}$

$\frac{\log n}{2}$	KBs
8	0.75
16	512

Small bitmaps

- Rank takes constant time on small bitmaps (a computer word)
- Same idea works for select, the possible answers for a block are either a position or “not present”
- Together with the tables for rank, that is enough for answering select in constant time for small bitmaps


Small bitmaps

- In practice, we can use processor instructions to replace the rank tables (this is called popcount).

```
inline cds_word popcount(cds_word x) {  
    if (unlikely(x == 0)) {  
        return 0;  
    }  
    return __builtin_popcountl(x);  
}
```


Rank

- We know how to solve for small bitmaps, so try reducing to that
- Lets start by storing some partial answers every s bits

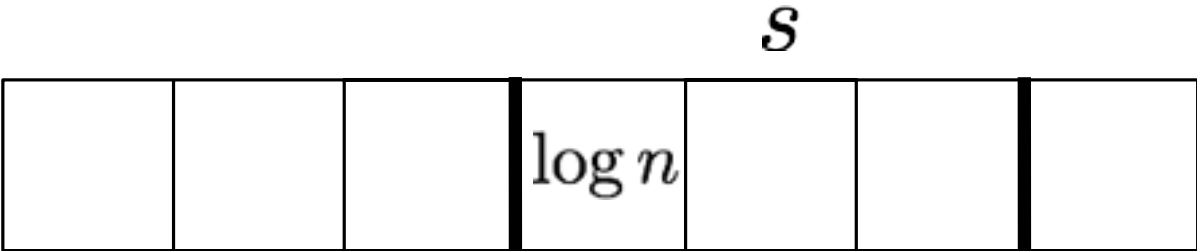


Rank

- We know how to solve for small bitmaps, so try reducing to that
- Lets start by storing some partial answers every s bits



Rank


- Sampling every s elements 

			$\log n$			
--	--	--	----------	--	--	--
- We can answer rank in $O\left(\frac{s}{\log n}\right)$ time using samples and the tables
- The way to improve further is to consider the blocks generated by the samples as independent problems


Rank

- We sample every $b \leq s$ bits, each sample requires $\log s$ bits
- We want to be left with blocks of size $\frac{\log n}{2}$
- We achieve this setting $b = \frac{\log n}{2}$ and $s = \log^2 n$


Rank




Rank




Rank




Rank



Rank



Rank



3 additions + popcount!

Rank

- Tables: $\frac{\sqrt{n} \log n \log \log n}{2}$
- Blocks: $\frac{2n \log \log n}{\log n}$
- Super blocks: $\frac{n}{\log n}$
- Bitmap: n
- **Total:** $n + O\left(\frac{n \log \log n}{\log n}\right) = n + o(n)$

Select

- We can try something similar to rank, but there is a catch: we cannot use fixed-sized blocks.

B =

10100101	001001011	0010101001
----------	-----------	------------

\longleftrightarrow
 $\log^2 n$ 1s

Select

- We know the answer every $\log^2 n$ Is and this generates blocks
- We split into two cases: sparse and dense blocks
- We store the answer for all possible arguments for sparse blocks, and recurse on the dense ones

Select

- Sparse blocks (length at least $\log^4 n$):
 - Each answer requires $\log n$ bits
 - The maximum space we will spend is:

$$\frac{n}{\log^4 n} \times \log^3 n = \frac{n}{\log n}$$

Select

- Dense blocks (length at most $\log^4 n$):
 - Split into blocks with $(\log \log n)^2$ ones
 - These sub-blocks are classified as sparse or dense
 - A sub-block is sparse if its length is at least

$$4 \times (\log \log n)^4$$

- Same idea as before, now the overhead is:

$$\frac{n}{4(\log \log n)^4} (\log \log n)^2 \times 4 \log \log n = \frac{n}{\log \log n}$$

Select

- Answering a query:
 - If the block is sparse, return the answer
 - Else go to the corresponding sub-block
 - If the sub-block is sparse, return the answer
 - Else it is not sparse, but it fits in a word

In practice...

- We only keep one level of blocks for rank
- Select is solved the following way:
 - Binary search which block contains the answer
 - Sequentially traverse the block to find the position
- There is another solution storing samples for select + the ones for rank

In practice...

```
Array *a = Array::Create(n, 1);  
...  
BitSequence *bs = new BitSequenceOneLevelRank(a, sample);  
cout << bs->Access(i) << endl;  
cout << bs->Rank0(i) << " " << bs->Rank1(i) << endl;  
cout << bs->Select0(i) << " " << bs->Select1(i) << endl;
```

Other representations

- Raman, Raman and Rao (constant time)

$$nH_o(B) + o(n)$$

- Okanohara and Sadakane (not constant time)


$$nH_o(B) + O(m)$$

- Patrascu (constant time)
 - Reduced the lower order term for compressed bitmaps

Huffman + bitmap

- To access position i , we just do


$\text{decode}(H[\text{select}(l, i) \dots \text{select}(l, i+1)-1])$



Huffman + bitmap

- To access position i , we just do

$\text{decode}(H[\text{select}(l, i) \dots \text{select}(l, i+1)-1])$



Going back to Huffman


- We can use a compressed bitmap instead
- We mark the beginning of each code in a bitmap of length
- There are n ones $nH_0(S) + n$
- The space for the whole representation is

$$nH_0(S) + o(nH_0(S))$$

Huffman

- A pointer-based representation requires $O(n \log n)$ bits


```
class Node {  
    void *data;  
    Node *left, *right;  
}
```



Huffman


- A pointer-based representation requires $O(n \log n)$ bits

```
class Node {  
    void *data;  
    Node *left, *right;  
}
```



Huffman

- A pointer-based representation requires $O(n \log n)$ bits




```
class Node {  
    void *data;  
    Node *left, *right;  
}
```

What if we want to know our parent?


Huffman

- Good example of a simple tree
- Every node has 2 children or is a leaf
- Can we represent the shape of the tree efficiently?




Simple representation

- We will traverse the tree DFS
- Every time we see an internal node, we write a 1
- Every time we see a leaf, we write a 0



Simple representation


- We will traverse the tree DFS
- Every time we see an internal node, we write a 1
- Every time we see a leaf, we write a 0



Simple representation

- We will traverse the tree DFS
- Every time we see an internal node, we write a 1
- Every time we see a leaf, we write a 0


1



Simple representation

- We will traverse the tree DFS
- Every time we see an internal node, we write a 1
- Every time we see a leaf, we write a 0


1



Simple representation

- We will traverse the tree DFS
- Every time we see an internal node, we write a 1
- Every time we see a leaf, we write a 0


1 1



Simple representation

- We will traverse the tree DFS
- Every time we see an internal node, we write a 1
- Every time we see a leaf, we write a 0


1 1



Simple representation

- We will traverse the tree DFS
- Every time we see an internal node, we write a 1
- Every time we see a leaf, we write a 0


1 1 0



Simple representation

- We will traverse the tree DFS
- Every time we see an internal node, we write a 1
- Every time we see a leaf, we write a 0


1 1 0



Simple representation

- We will traverse the tree DFS
- Every time we see an internal node, we write a 1
- Every time we see a leaf, we write a 0


1 1 0 1



Simple representation

- We will traverse the tree DFS
- Every time we see an internal node, we write a 1
- Every time we see a leaf, we write a 0


1 1 0 1 0



Simple representation

- We will traverse the tree DFS
- Every time we see an internal node, we write a 1
- Every time we see a leaf, we write a 0


1 1 0 1 0 0




Simple representation

- We will traverse the tree DFS
- Every time we see an internal node, we write a 1
- Every time we see a leaf, we write a 0

1 1 0 1 0 0 0




Can we navigate this?




1101000

Can we navigate this?



1101000
▲


Can we navigate this?



1 1 0 1 0 0 0
 ▲

Who's my parent?

Can we navigate this?




1101000

Who's my parent?


Another try

- Let's give our tree another try
- Now using BFS



Another try


- Let's give our tree another try
- Now using BFS



Another try

- Let's give our tree another try
- Now using BFS


I



Another try

- Let's give our tree another try
- Now using BFS


I



Another try

- Let's give our tree another try
- Now using BFS


||



Another try

- Let's give our tree another try
- Now using BFS


||



Another try

- Let's give our tree another try
- Now using BFS


110



Another try

- Let's give our tree another try
- Now using BFS


1100



Another try

- Let's give our tree another try
- Now using BFS


1 1 0 0 1



Another try

- Let's give our tree another try
- Now using BFS


1 1 0 0 1 0




Another try

- Let's give our tree another try
- Now using BFS

1 1 0 0 1 0 0




Another try (part 2)




Another try (part 2)

I
II
I00I
0000




Another try (part 2)

|
| |
| 0 0 |
 ▲
0 0 0 0




Another try (part 2)

1
1 1
↑ 001
0000




Another try (part 2)

I
II
I00I
0000




Another try (part 2)

|
| |
| 0 0 |
0 0 0 0




Another try (part 2)


|
| |
| 00 |
0000



LOUDS




LOUDS




LOUDS

1110




LOUDS

1110




LOUDS

111010




LOUDS

111010110110001100000



LOUDS


b c d	e	f g	i j			k l				
1 1 1 0	1 0	1 1 0	1 1 0 0 0	0 0	1 1 0 0 0 0 0					
a	b	c	d	e f	g	i j k l				



LOUDS

b c d	e	f g	i j			k l				
1 1 1 0	1 0	1 1 0	1 1 0	0 0	0 0	1 1 0	0 0	0 0	0 0	0 0
a	b	c	d	e	f	g	i	j	k	l

Where is the k-th
node in BFS order?




LOUDS

b c d	e	f g	i j			k l				
1 1 1 0	1 0	1 1 0	1 1 0	0 0	0 0	1 1 0	0 0	0 0	0 0	0 0
a	b	c	d	e	f	g	i	j	k	l

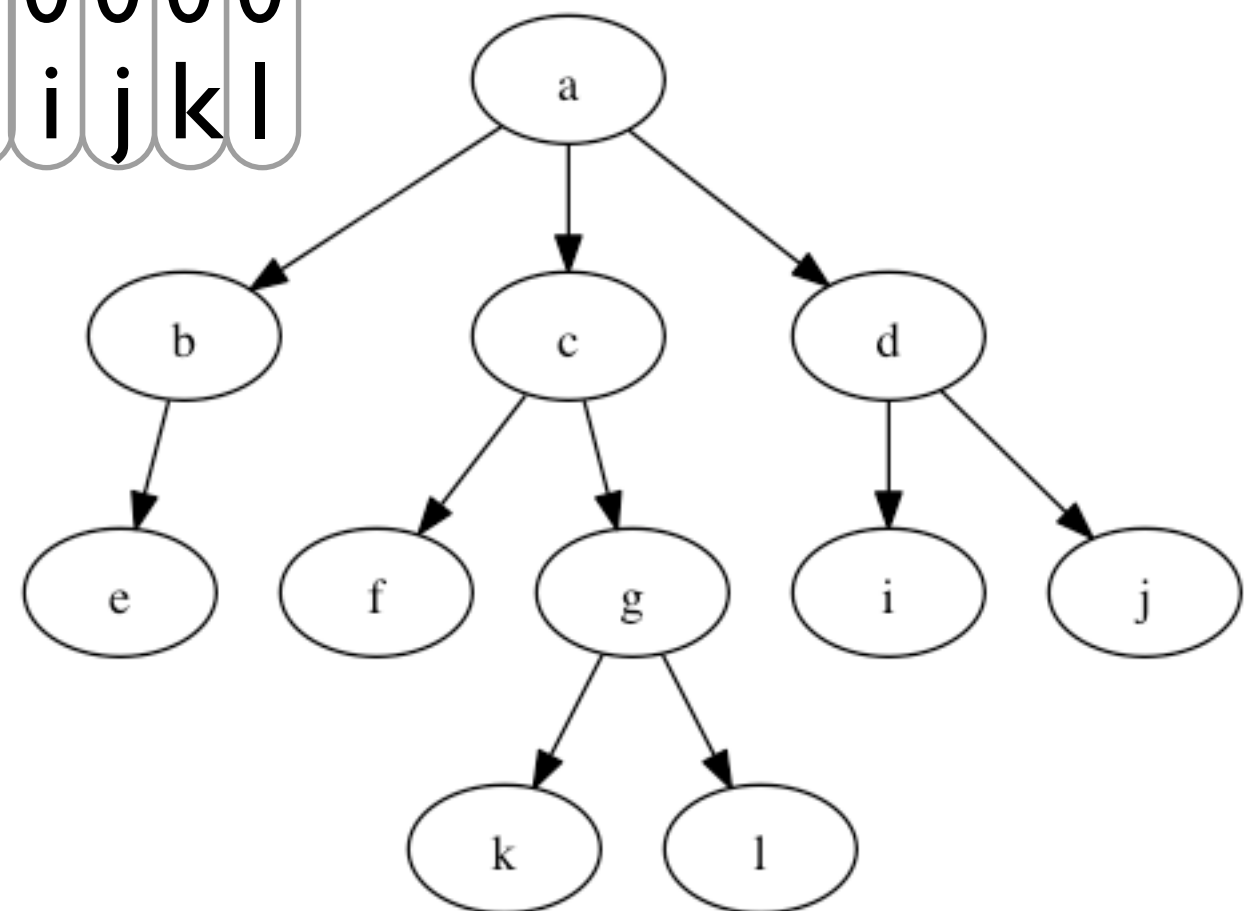
Where is the k-th
node in BFS order?

$\text{select}(0, k)$



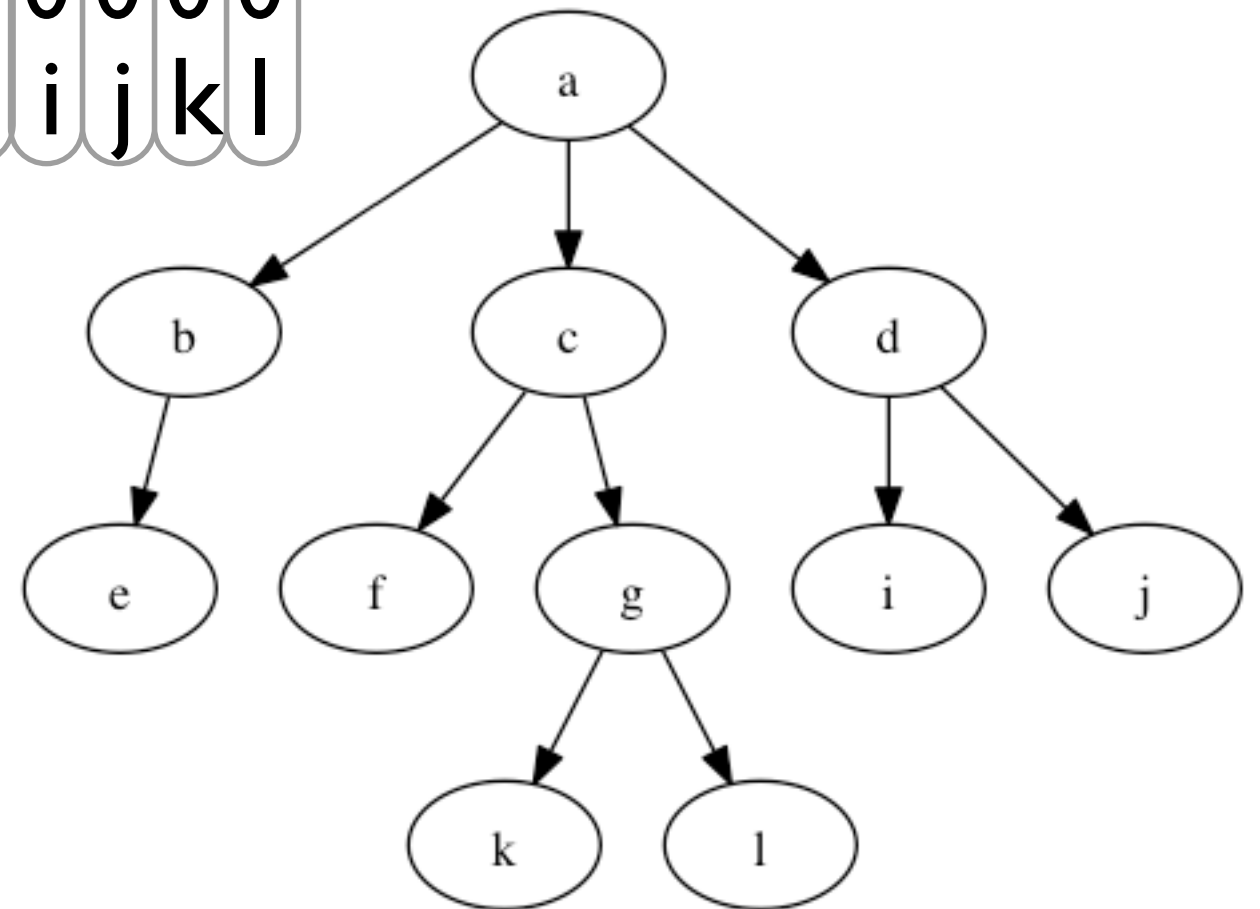
LOUDS

b c d	e	f g	i j			k l				
1 1 1 0	1 0	1 1 0	1 1 0 0 0	0 0	1 1 0 0 0 0 0					
a	b	c	d	e	f	g	i	j	k	l

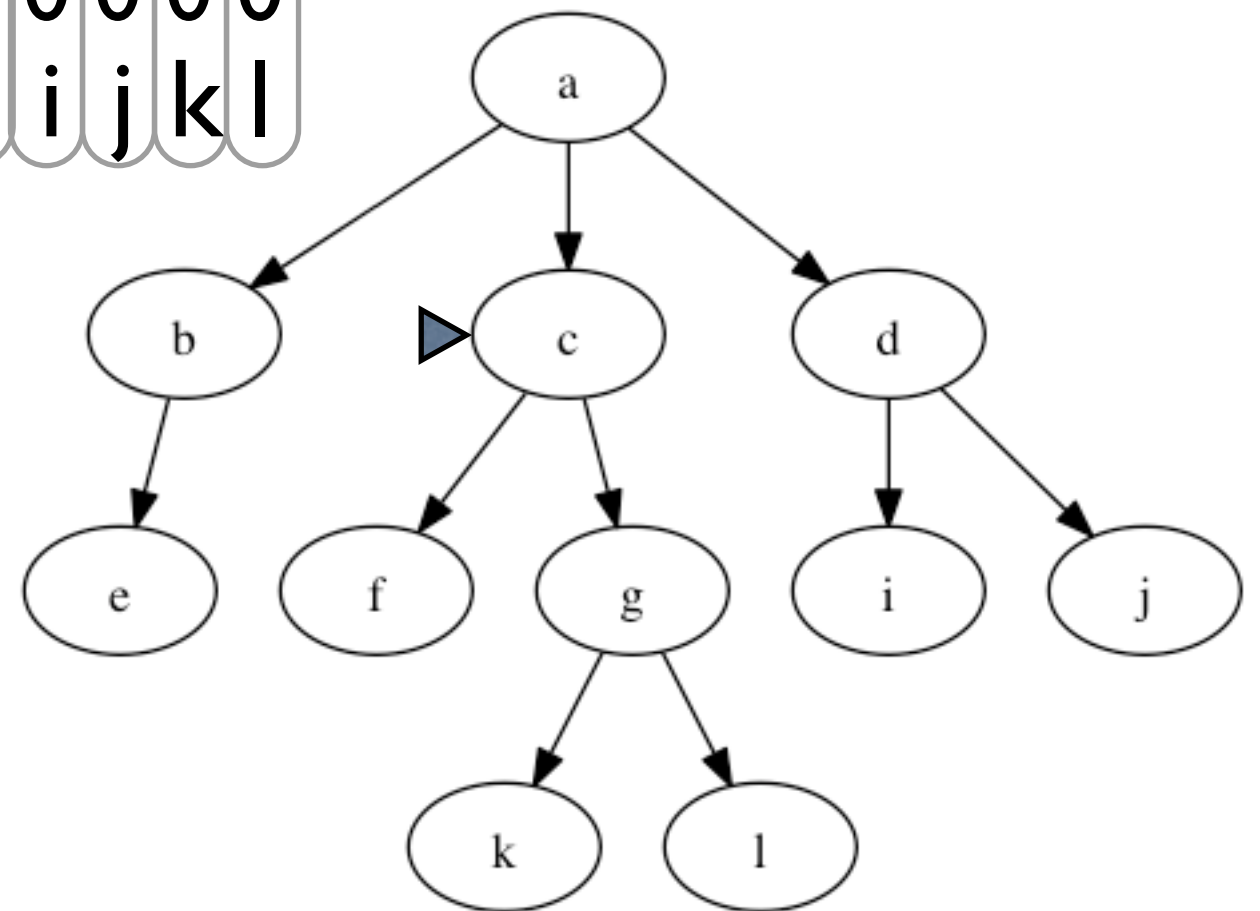



LOUDS (child)



b c d	e	f g	i j			k l				
1 1 1 0	1 0	1 1 0	1 1 0 0 0	0	0	1 1 0 0 0 0 0	0	0	0	0
a	b	c	d	e	f	g	i	j	k	l




LOUDS (child)




LOUDS (child)




LOUDS (child)



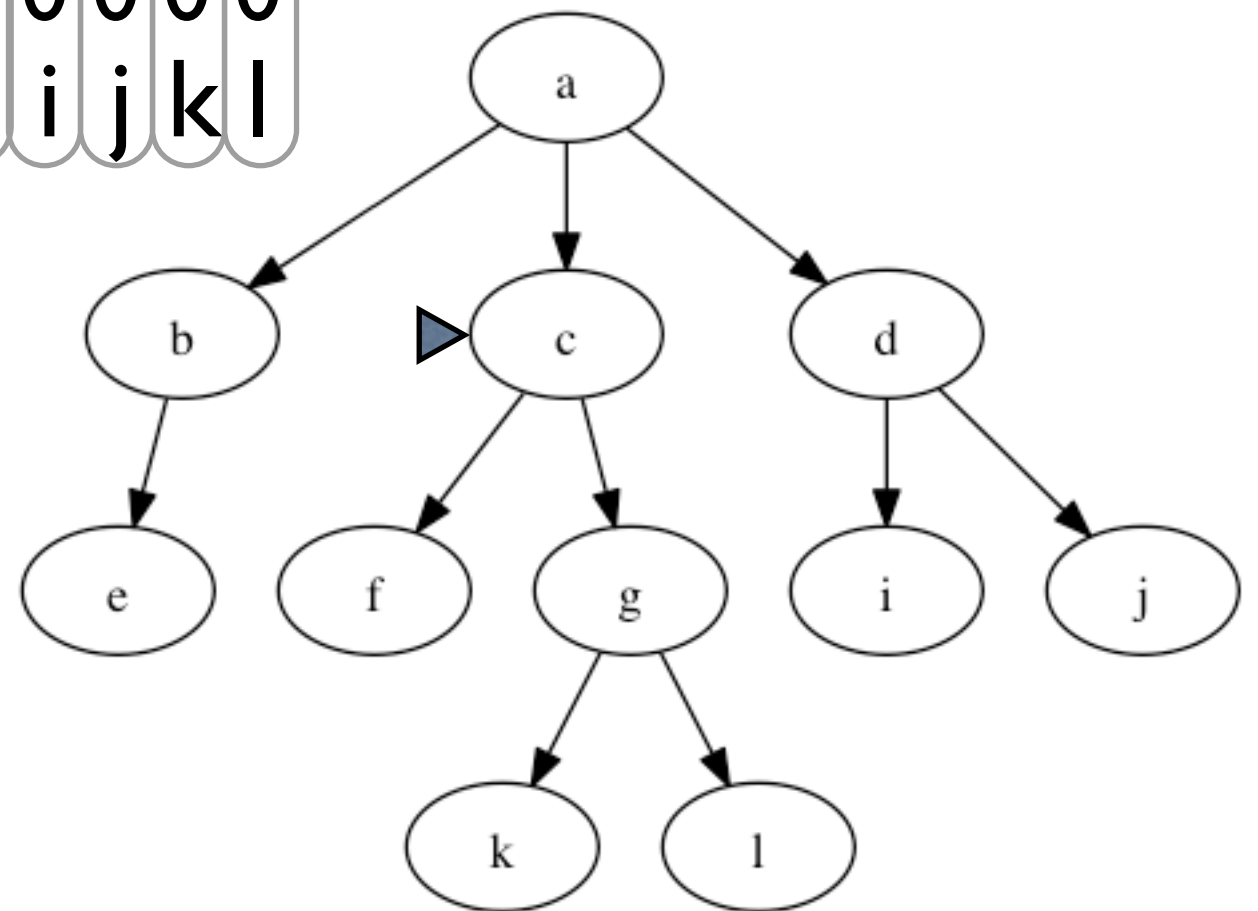
f is the 6th node



LOUDS (child)

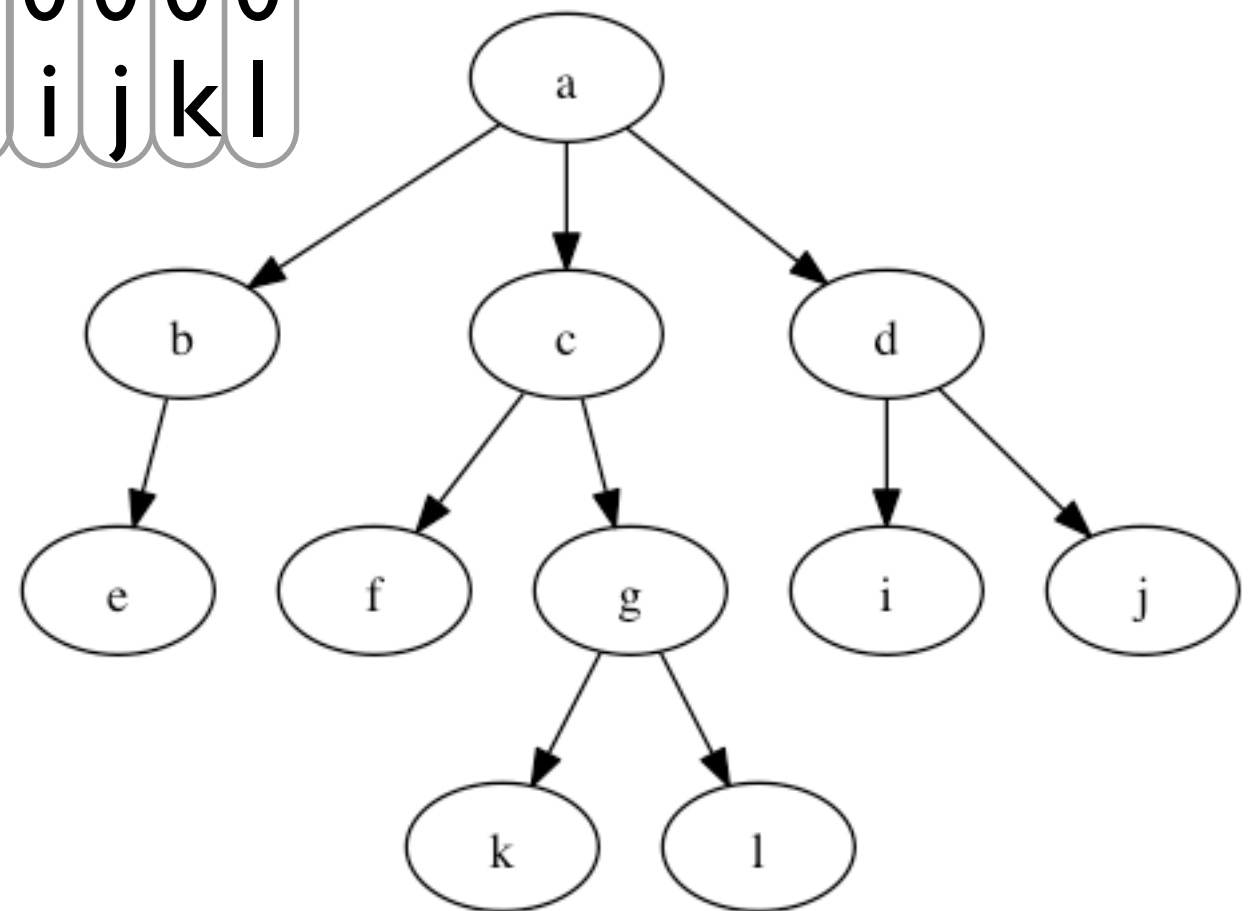


f is the 6th node
use `select(0, 6)`

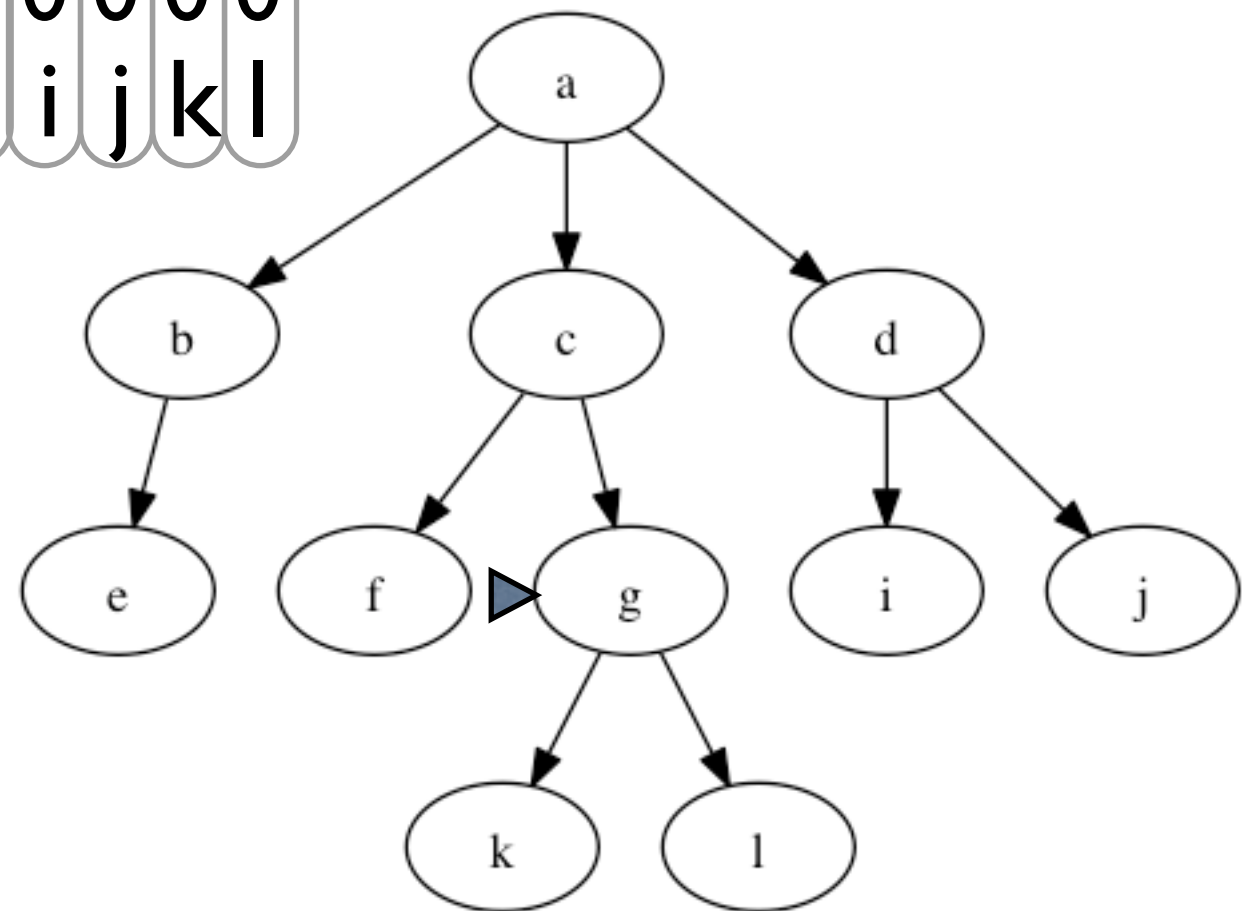



LOUDS (parent)

b c d	e	f g	i j			k l				
1 1 1 0	1 0	1 1 0	1 1 0 0 0	0 0	1 1 0 0 0 0 0					
a	b	c	d	e	f	g	i	j	k	l

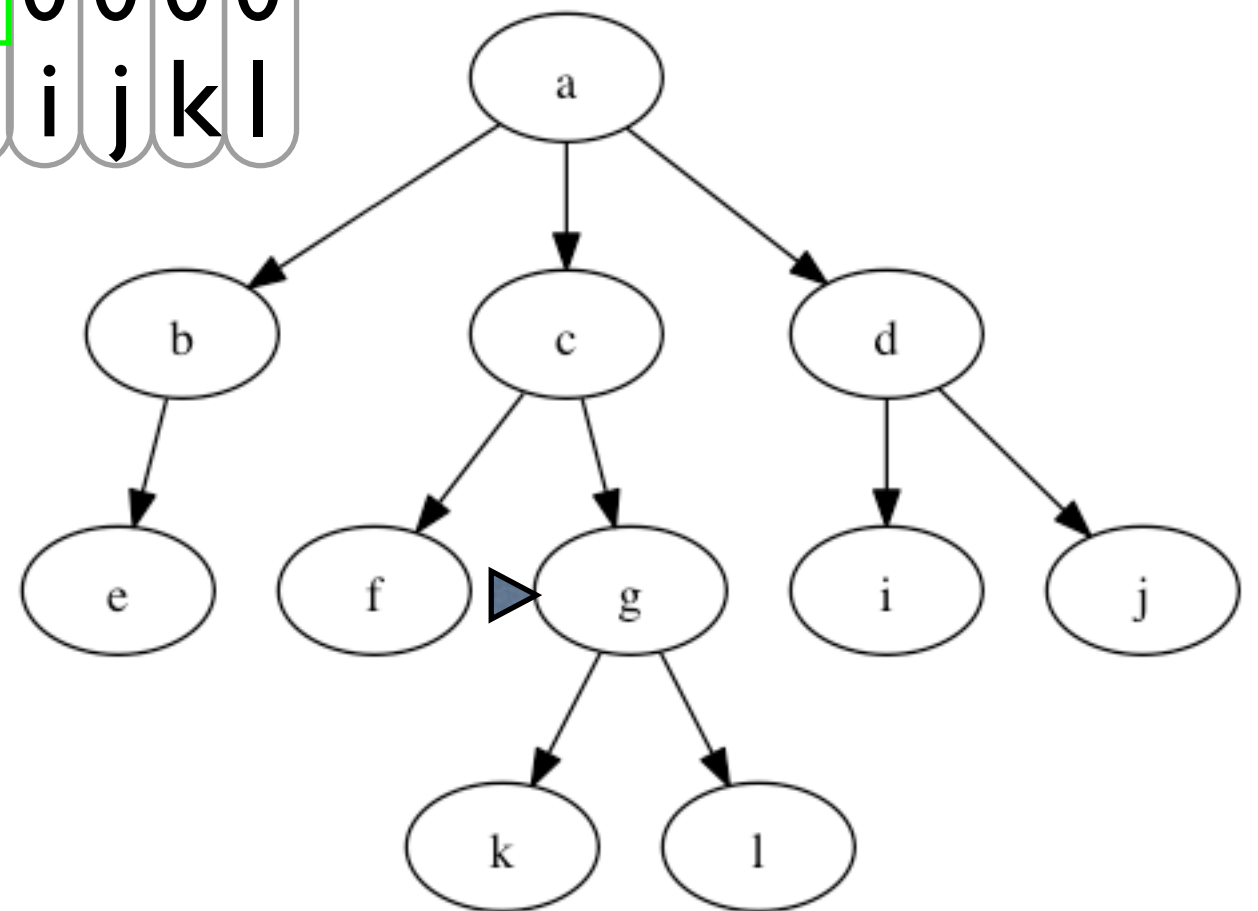


LOUDS (parent)



LOUDS (parent)

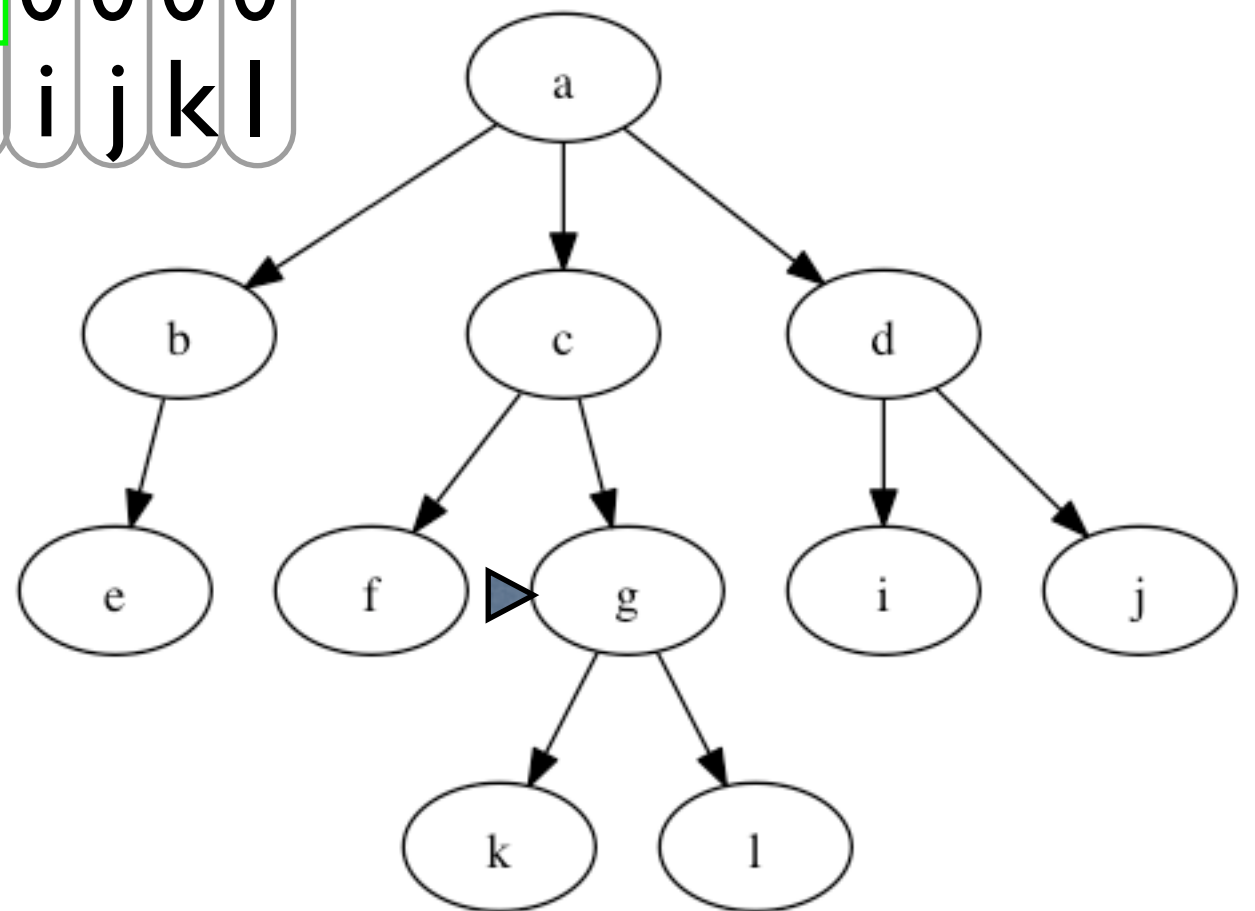
b	c	d	e	f	g	i	j			k	l				
1	1	1	0	1	0	1	1	0	0	0	1	1	0	0	0
a	b	c	d	e	f	g	i	j	k	l					



LOUDS (parent)

b c d	e	f g	i j			k l				
1 1 1 0	1 0	1 1 0	1 1 0	0 0	0 0	1 1 0	0 0	0 0	0 0	0 0
a	b	c	d	e	f	g	i	j	k	l


g is the 7th node



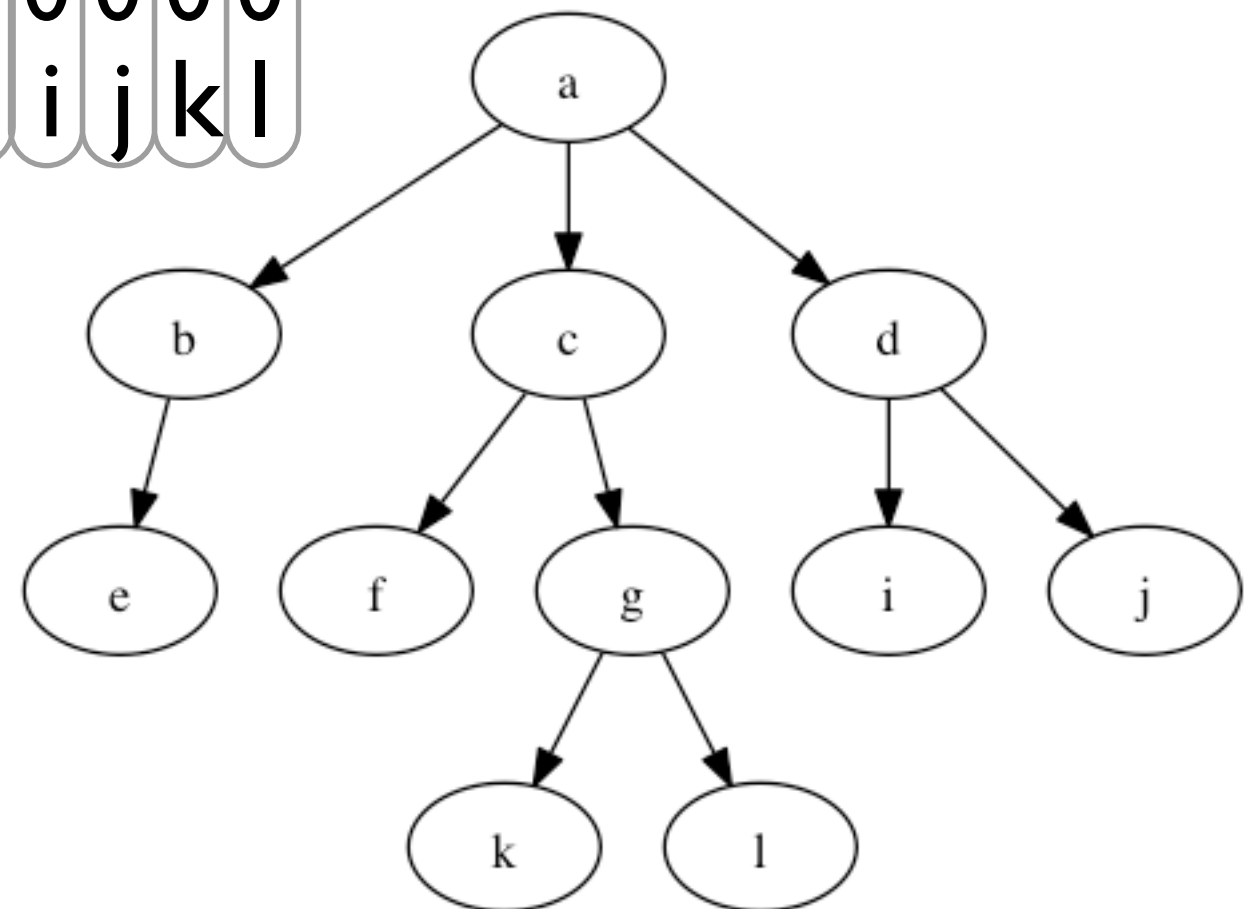

LOUDS (parent)

b c d	e	f g	i j			k l				
1 1 1 0	1 0	1 1 0	1 1 0	0 0	0 0	1 1 0	0 0	0 0	0 0	0 0
a	b	c	d	e	f	g	i	j	k	l


g is the 7th node
use $\text{select}(1, 7 - 1)$




LOUDS (others)



LOUDS (others)



Left/right sibling




LOUDS (others)

b c d	e	f g	i j			k l				
1 1 1 0	1 0	1 1 0	1 1 0	0 0	0 0	1 1 0	0 0	0 0	0 0	0 0
a	b	c	d	e	f	g	i	j	k	l

Left/right sibling

Degree



LOUDS (others)

b	c	d		e		f	g		i	j				k	l											
1	1	1	0	1	0	1	1	0	1	1	0	0	0	1	1	0	0	0	0							
a				b		c			d			e		f		g			i		j		k		l	

```
Tree *createTree() {
    vector<cds_word> v(21);
    v[0] = v[1] = v[2] = v[4] = v[6] = v[7] = 1;
    v[9] = v[10] = v[14] = v[15] = 1;
    Array *a = Array::Create(v);
    return new TreeLouds(new BitSequenceOneLevelRank(a, 20));
}
```

```

...
Tree * t = createTree();
assert(t->Parent(12) == 5);
assert(t->Child(5, 0) == 12);
assert(t->Degree(16) == 2);
... t->NextSibling(12)
... t->PrevSibling(12)
...


```

Trees

- BP: Geary, Raman, Raman & Rao
- DFUDS: Benoit et al. '05
- FF: Sadakane & Navarro '10
- Partitioning: Farzan & Munro '09

Trees

- There are some amazing results related to trees




Trees

- There are some amazing results related to trees




Trees

- There are some amazing results related to trees



Trees

- There are some amazing results related to trees



$2n + o(n)$ bits and constant time!

Sequences

Sequences

- $\text{rank}(a, i)$: counts how many times does a occur up to position i
- $\text{select}(a, j)$: finds the j -th occurrence of a
- $\text{access}(i)$: retrieves the i -th element

Sequences

- $\text{rank}(a, i)$: counts how many times does a occur up to position i
- $\text{select}(a, j)$: finds the j -th occurrence of a
- $\text{access}(i)$: retrieves the i -th element

alabaralalabarda


$$\text{rank}(r, 11) = 1$$

Wavelet Trees

- The best known structure for solving rank, select, and access on sequences
- Has many other applications we will not have time to cover
- Wavelet trees are awesome!


Wavelet Trees

access(8)




Wavelet Trees

access(8)




Wavelet Trees

access(8)




Wavelet Trees


access(8)




Wavelet Trees




Wavelet Trees



Wavelet Trees




Wavelet Trees




Wavelet Trees

select('E', 3)




Wavelet Trees

select('E', 3)




Wavelet Trees

select('E', 3)



Wavelet Trees

select('E', 3)



Wavelet Trees

- Each level has n bits
- There are $\lceil \log \sigma \rceil$ levels
- All queries cost $O(\log \sigma)$ time
- The total space is $n \log \sigma + o(n \log \sigma)$ bits

Wavelet Trees

- We can also compress the sequence, changing the shape of the tree.
- Any encoding works.
- With Huffman shape we get

$$n(H_0(S) + 1) + o(n(H_0 + 1))$$

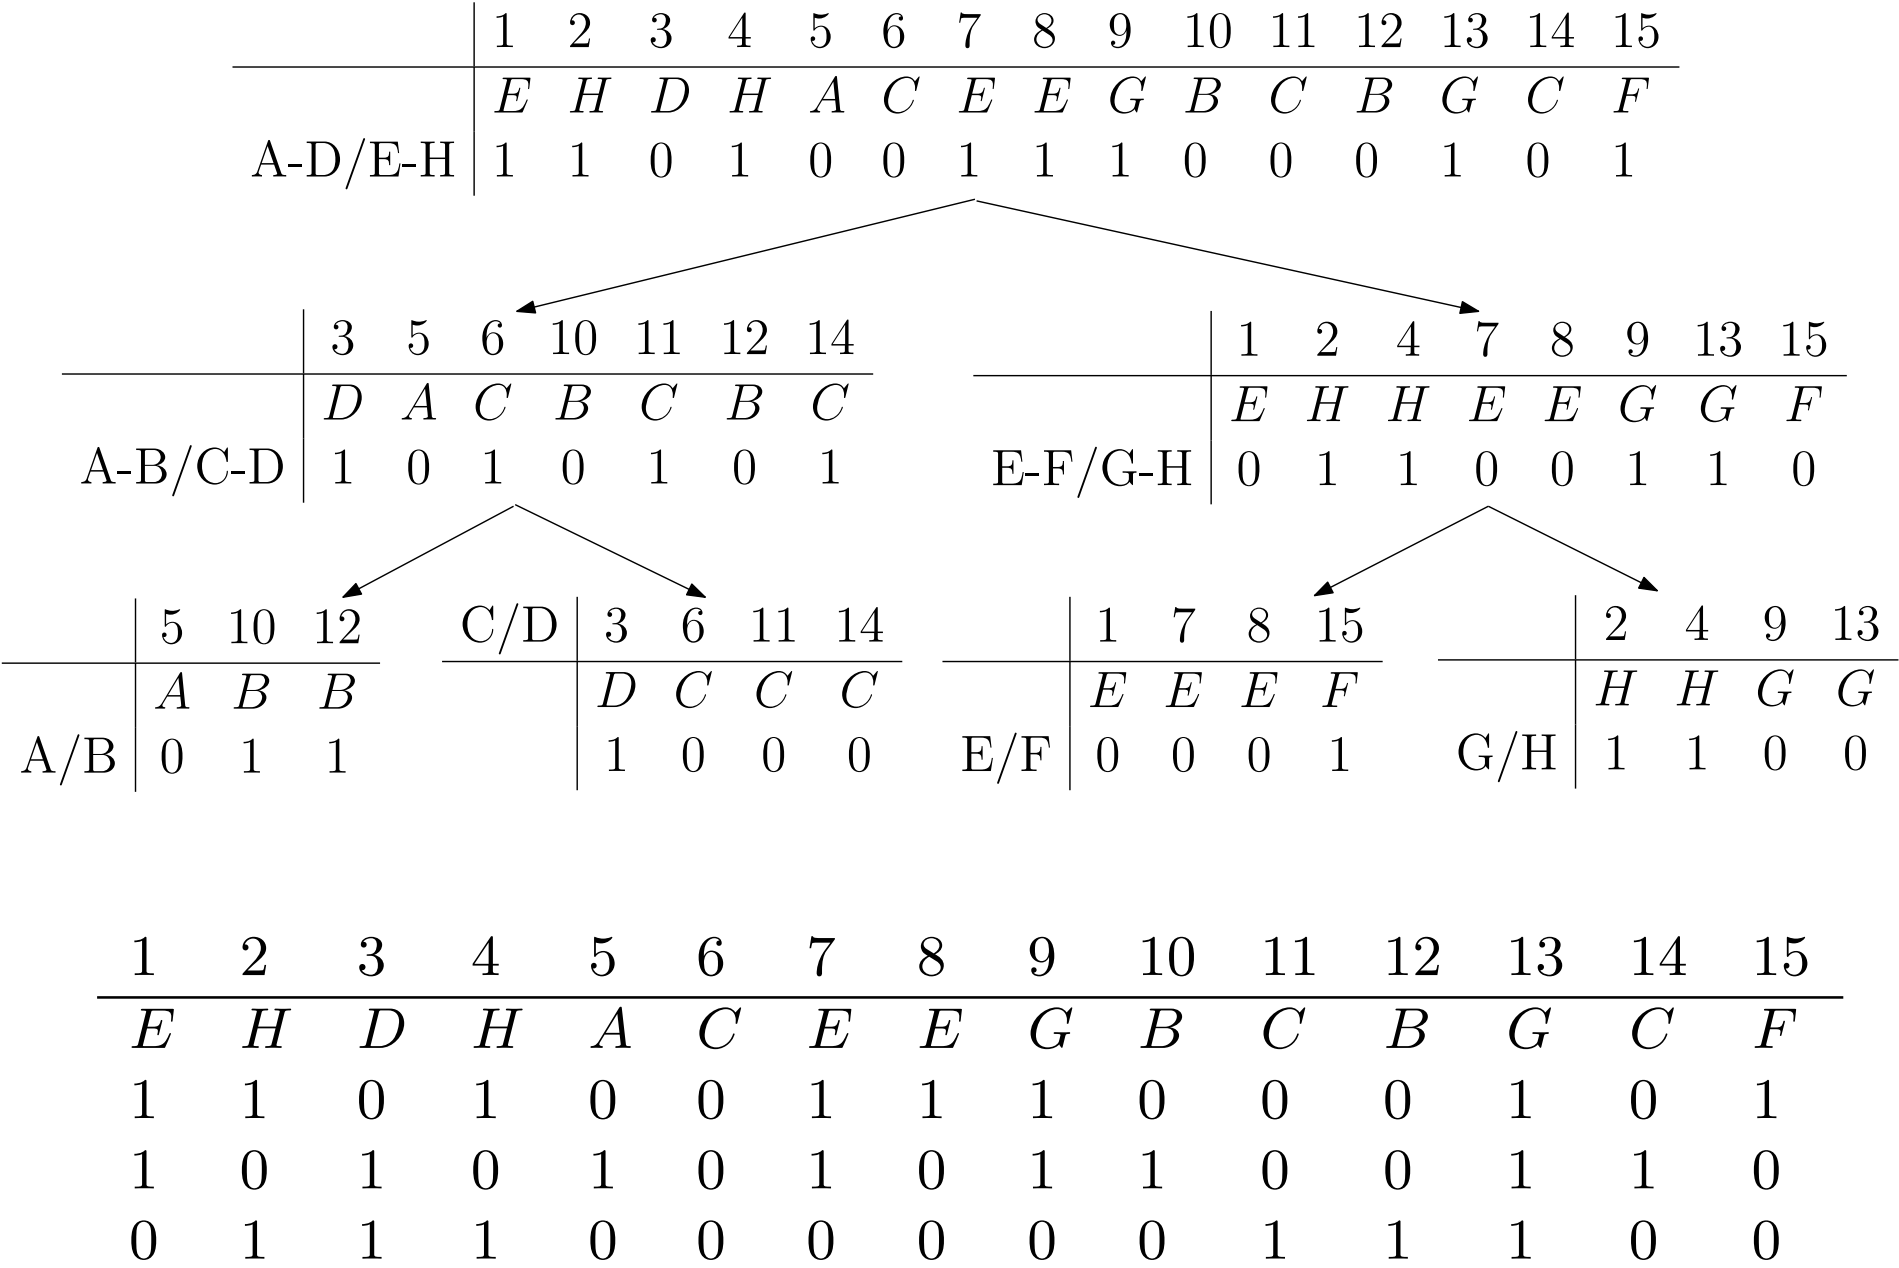
Wavelet Trees

- What about compressing the bitmaps?
- If we use RRR, we also get close to H_0
- It works really well on sequences with runs, like the BWT.

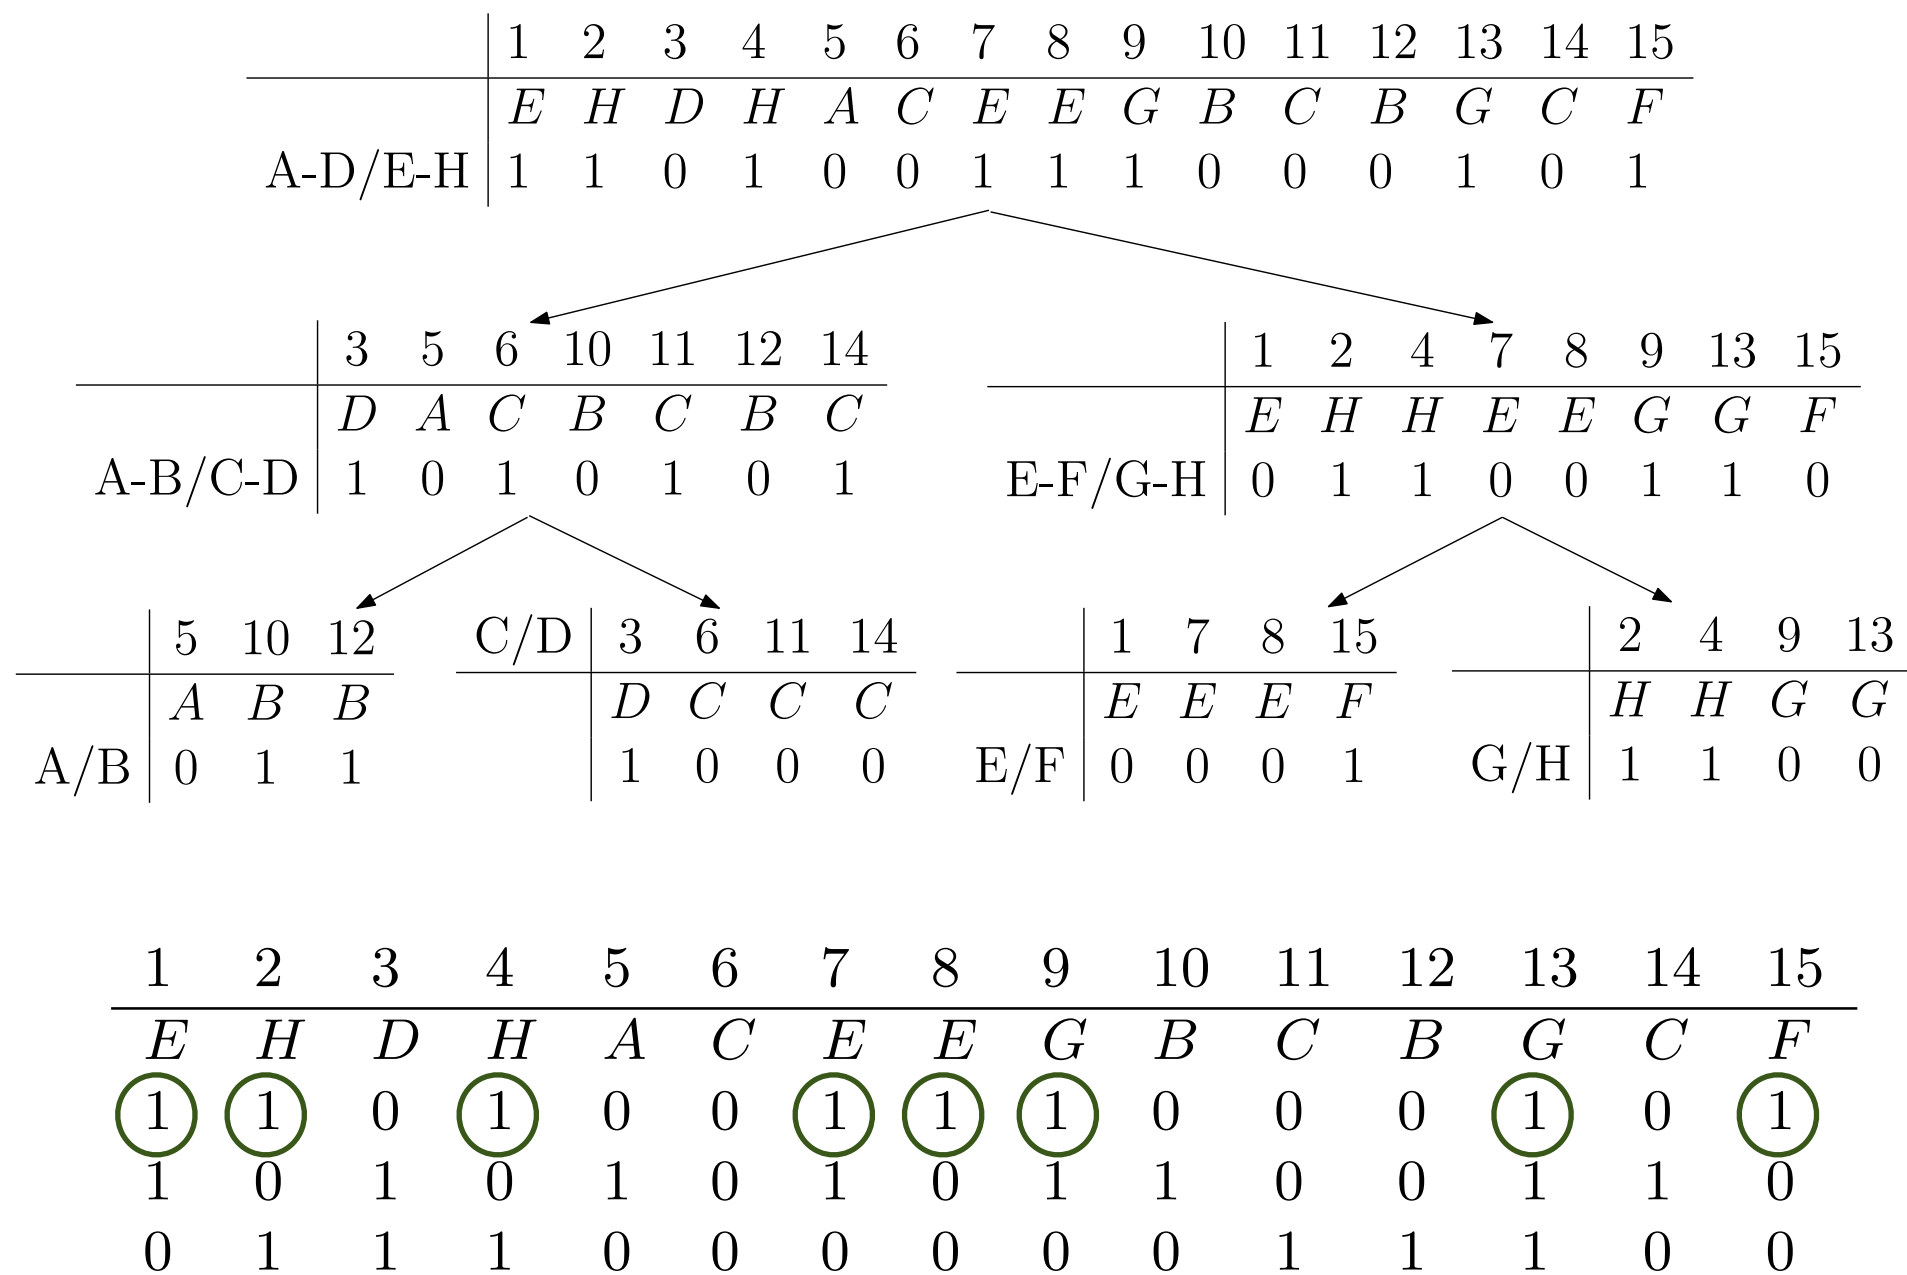
Wavelet Trees

- One problem: the tree
- We are spending $O(\sigma \log \sigma)$ bits on pointers!


Wavelet Trees



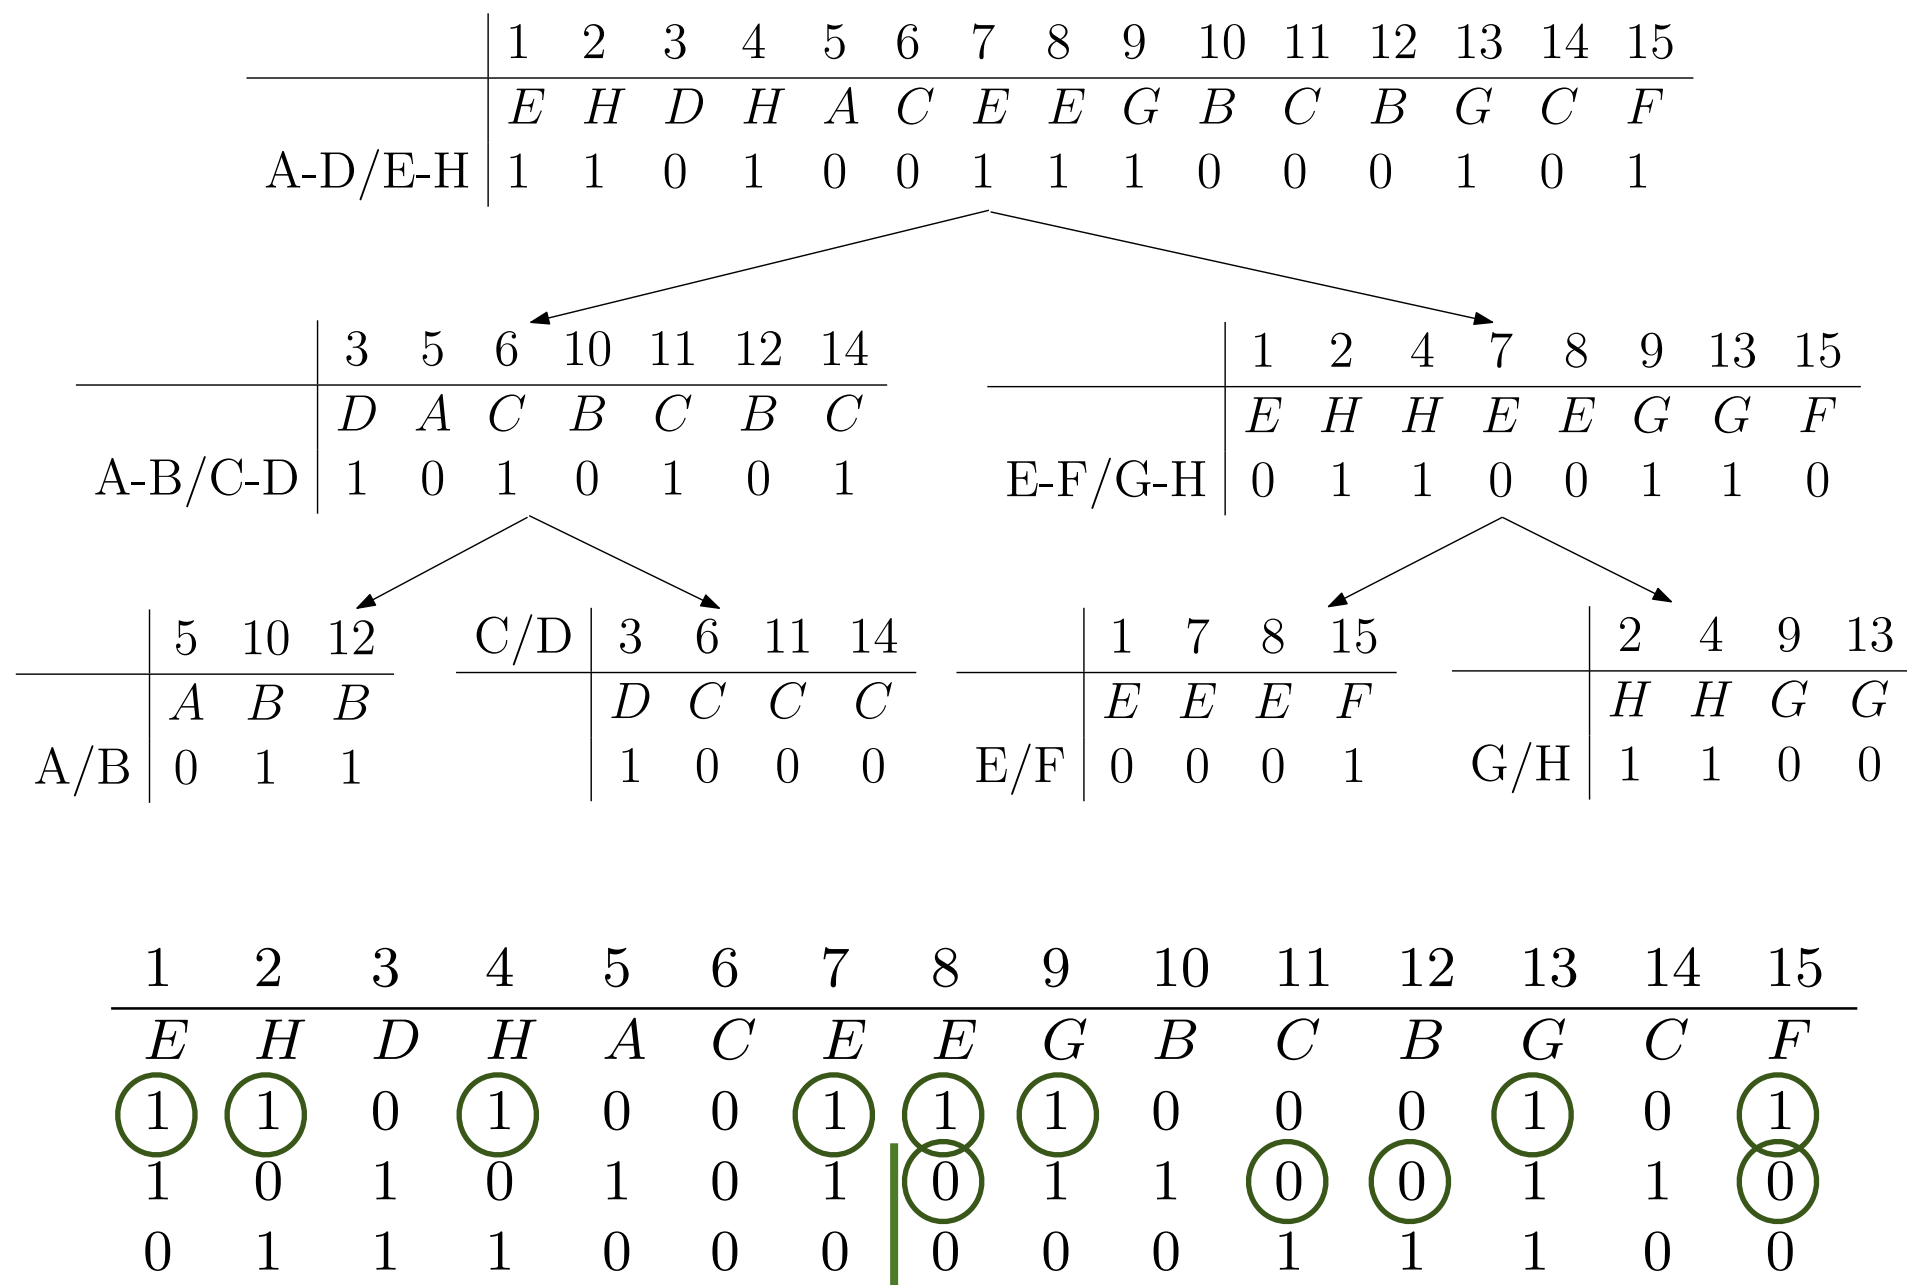
Wavelet Trees



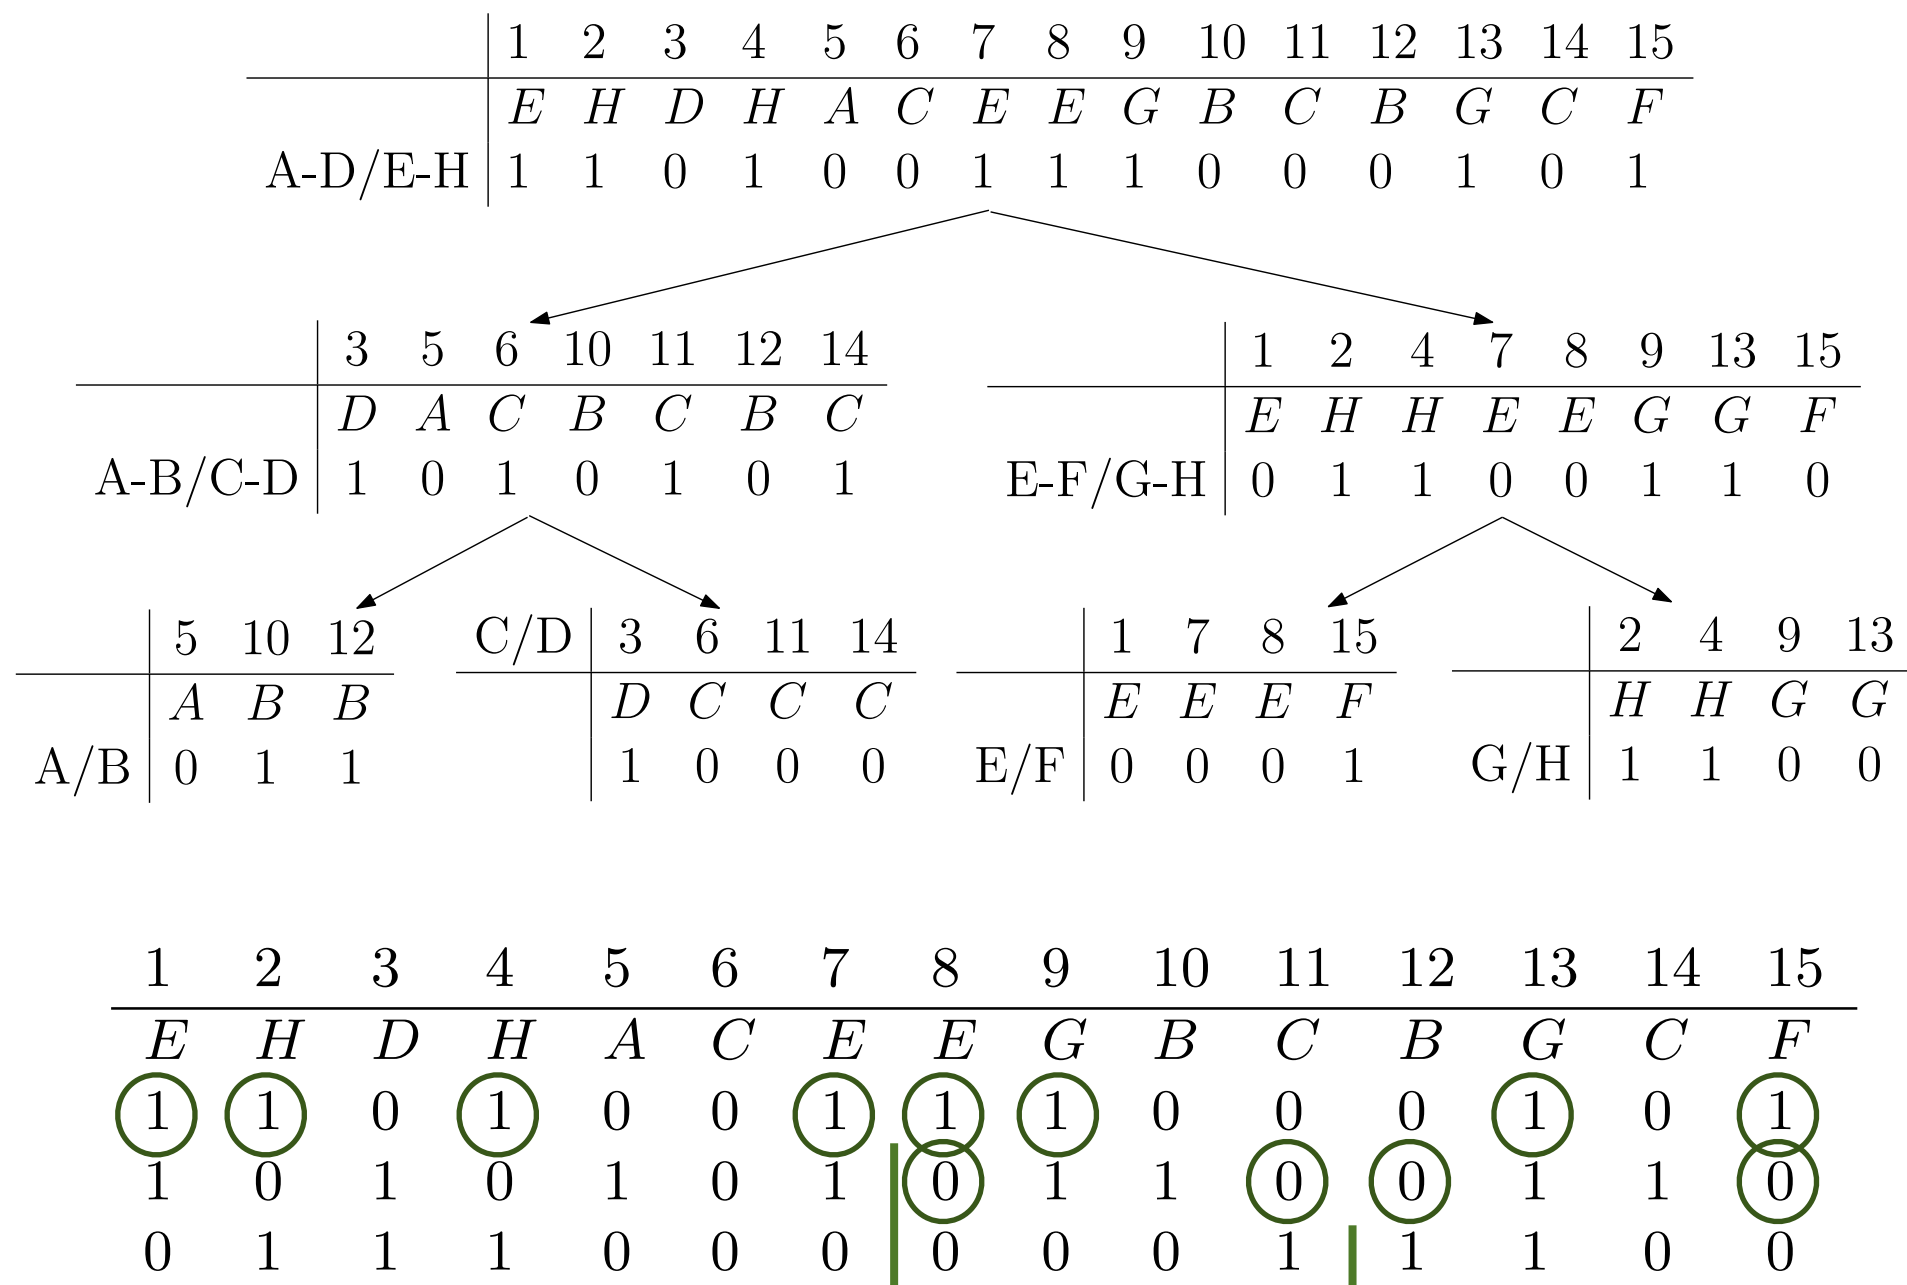
Wavelet Trees



Wavelet Trees



Wavelet Trees



Wavelet Trees

- We now need $n \log \sigma(1 + o(1))$ bits
- We don't waste $O(\sigma \log \sigma)$ bits in pointers
- To move from one level to another we perform 2 rank queries

Wavelet Trees

```
Array *array = CreateRandomSequence(len, sigma);  
WaveletTreeNoPtrs *seq = new WaveletTreeNoPtrs(array,  
                                                new BitSequenceBuilderOneLevelRank(20),  
                                                new MapperNone());  
  
seq->Rank(65, 10);  
seq->Select('A', 5);  
seq->Access(10);  
...
```

Wavelet Matrix

- Can we reduce the number of operations when we move from one level to the other?

Wavelet Matrix

- Can we reduce the number of operations when we move from one level to the other?

YES! Go to our talk on Wednesday :-)

Wavelet Trees

- Can we give Huffman shape to the wavelet tree without pointers?

Wavelet Trees

- Can we give Huffman shape to the wavelet tree without pointers?

YES! Go to Alberto's talk tomorrow :-)

Permutations

Representing a permutation of size n
requires $n \log n$ bits

$$\pi = [3, 4, 9, 2, 1, 5, 7, 6, 8]$$

Permutations

Representing a permutation of size n
requires $n \log n$ bits

$$\pi = [3, 4, 9, 2, 1, 5, 7, 6, 8]$$

What if we want to compute $\pi^{-1}(i)$

Permutations

$$\pi = [3, 4, 9, 2, 1, 5, 7, 6, 8]$$

Permutations

$$\pi = [3, 4, 9, 2, 1, 5, 7, 6, 8]$$

$\pi(i)$ is easy

Permutations

$$\pi = [3, 4, 9, 2, 1, 5, 7, 6, 8]$$

$\pi(i)$ is easy

$\pi^{-1}(i)$ is also easy, if we spend $n \log n$ extra bits

Permutations

$$\pi = [3, 4, 9, 2, 1, 5, 7, 6, 8]$$


$\pi(i)$ is easy

$\pi^{-1}(i)$ is also easy, if we spend $n \log n$ extra bits

Wavelet trees solve this in $O(\log n)$


Permutations

$$\pi = [3, 4, 9, 2, 1, 5, 7, 6, 8]$$



Permutations


$$\pi = [3, 4, 9, 2, 1, 5, 7, 6, 8]$$



Permutations

$$\pi = [3, 4, 9, 2, 1, 5, 7, 6, 8]$$


$$\pi^{-1}(8)$$



Permutations

$$\pi = [3, 4, 9, 2, 1, 5, 7, \underset{\blacktriangle}{6}, 8]$$


$$\pi^{-1}(8)$$



Permutations

$$\pi = [3, 4, 9, 2, 1, \underset{\blacktriangle}{5}, 7, \underset{\blacktriangle}{6}, 8]$$


$$\pi^{-1}(8)$$



Permutations

$$\pi = [3, 4, 9, 2, \underset{\blacktriangle}{1}, \underset{\blacktriangle}{5}, 7, \underset{\blacktriangle}{6}, 8]$$


$$\pi^{-1}(8)$$



Permutations

$$\pi = [3, 4, 9, 2, 1, 5, 7, 6, 8]$$


$$\pi^{-1}(8)$$



Permutations

$$\pi = [3, 4, 9, 2, 1, 5, 7, \underset{\blacktriangle}{6}, 8]$$


$$\pi^{-1}(8)$$



Permutations

$$\pi = [3, 4, 9, 2, 1, 5, 7, \underset{\blacktriangle}{6}, 8]$$

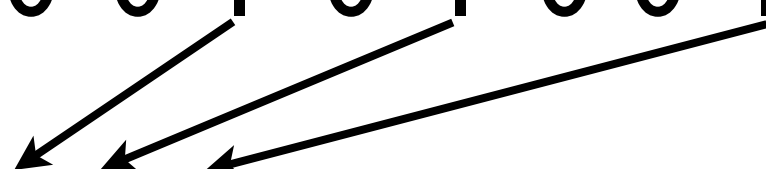
$$\pi^{-1}(8)$$




Permutations

$$\pi = [3, 4, 9, 2, 1, 5, 7, 6, 8]$$


0 0 | 0 | 0 0 | 0

$$P = [5, 8, 3]$$



Permutations




Permutations




Permutations



Permutations



Permutations



Permutations

$$\begin{array}{cccccccccc}
 & & \blacktriangledown & & & & \blacktriangledown & & \blacktriangledown & \\
 \pi = [3, 4, 9, 2, 1, 5, 7, 6, 8] & & & & & & & & & n \log n + O(n) \\
 0 \ 0 \ | \ 0 \ | \ 0 \ 0 \ | \ 0 & & & & & & & & & n + o(n) \\
 & \swarrow & \swarrow & \swarrow & & & & & & \\
 P = [5, 8, 3] & & & & & & & & & \leq n/t \log n + O(n)
 \end{array}$$

Permutations

$$\begin{array}{rcc}
 \pi = [3, 4, 9, 2, 1, 5, 7, 6, 8] & \begin{array}{c} \blacktriangledown \quad \quad \blacktriangledown \quad \quad \blacktriangledown \\ 0 \quad 0 \quad | \quad 0 \quad | \quad 0 \quad 0 \quad | \quad 0 \end{array} & \begin{array}{l} n \log n + O(n) \\ n + o(n) \end{array} \\
 P = [5, 8, 3] & \begin{array}{c} \swarrow \quad \swarrow \quad \swarrow \\ \leq n/t \log n + O(n) \end{array} & \\
 \hline
 n \log n + \frac{n \log n}{t} + O(n)
 \end{array}$$

Permutations

- We can trade space for time in computing π^{-1}
- For $t = \log \log n$ we get $n \log n + o(n \log n)$ bits

Permutations

```
cds_word a[] = {1,2,3,4,5,6,7,8,9,0};
Array *perm_a = Array::Create((cds_word*)a, 0, 9);
PermutationMRRR *perm = new PermutationMRRR(perm_a, 3);
for (cds_word i = 0; i < perm->GetLength(); i++) {
    cds_word expected = (i+1) % 10;
    ASSERT_EQ(expected, perm->Access(i));
    expected = (i + 10 - 1) % 10;
    ASSERT_EQ(expected, perm->Reverse(i));
}
```

GMR

- Last time we didn't know how to represent permutations
- Lets focus on sequences of length $O(\sigma)$

GMR

$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

	0	0		0	0		0
a			b			c	

GMR

$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

	0	0		0	0		0
a			b			c	

access(5)

GMR

$$S = [abca\textcolor{green}{b}bca]$$

$$\Pi = \begin{array}{ccccccc} [1, 4, 8, 2, 5, 6, 3, 7] \\ | \quad 0 \quad 0 \quad | \quad 0 \quad 0 \quad | \quad 0 \\ a \qquad \qquad b \qquad \qquad c \end{array}$$

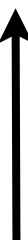
access(5)

GMR

$$S = [abca\textcolor{green}{b}bca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

	0	0		0	0		0
a			b			c	

$$\pi^{-1}(5)$$



access(5)

GMR

$$S = [abca\textcolor{green}{b}bca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

$\textcolor{green}{1}$	$\textcolor{green}{0}$	$\textcolor{green}{0}$	$\textcolor{green}{1}$	$\textcolor{green}{0}$	0	1	0
a			b			c	

$$\pi^{-1}(5)$$



access(5)

GMR

$$S = [abca\textcolor{green}{b}bca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

$\textcolor{green}{ }$	0	0	$\textcolor{green}{ }$	0	0	$\textcolor{green}{ }$	0
a			b			c	

$\pi^{-1}(5)$ 

$$\text{access}(5) = b$$


GMR

$$S = [abca\textcolor{green}{b}bca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

1	0	0	1	0	0	1	0
a			b			c	

$$\text{access}(5) = b$$

$$\pi^{-1}(5)$$


$$O(\log \log \sigma) \text{ time}$$

GMR

$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

	0	0		0	0		0
a			b			c	

select(b, 3)

GMR

$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

	0	0		0	0		0
a			b			c	

select(b, 3)

select(1, 2)

GMR

$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

	0	0		0	0		0
a			b			c	

select(b, 3)

select(1, 2)



GMR

$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

1	0	0	1	0	0	1	0
a			b			c	

select(b, 3)

select(1, 2)



GMR

$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

1	0	0	1	0	0	1	0
a			b			c	

select(b, 3)

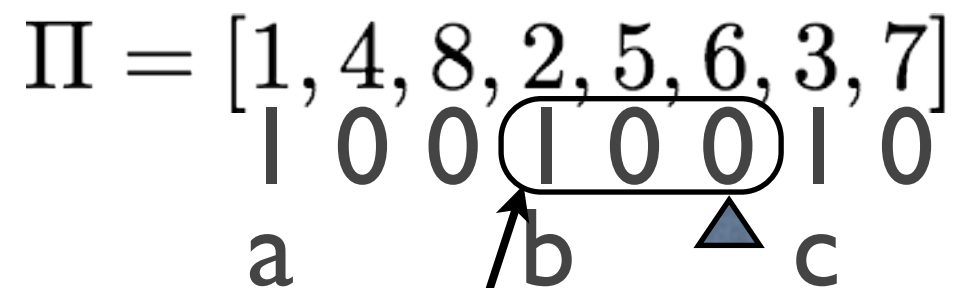
select(l, 2)

GMR

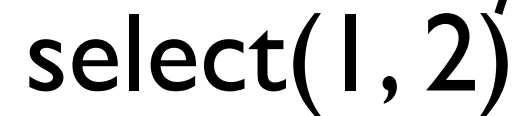
$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

	0	0		0	0		0
a			b			c	



$$\text{select}(b, 3) = 6$$

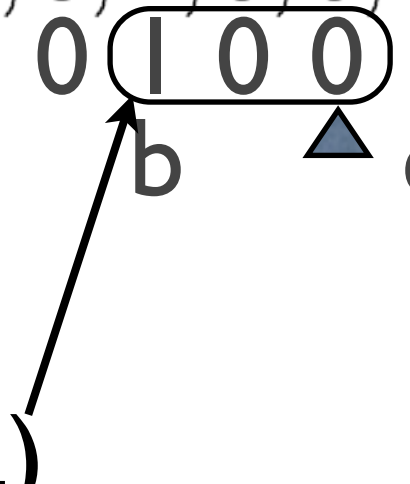
$$\text{select}(1, 2)$$


GMR

$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

	0	0		0	0		0
a			b			c	



$$\text{select}(b, 3) = 6$$

$$\text{select}(l, 2)$$

$$O(1) \text{ time}$$

GMR

$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

	0	0		0	0		0
a			b			c	

rank(b, 5)

GMR

$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

	0	0		0	0		0
a			b			c	

rank(b, 5)

select(1, 2)

GMR


$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

	0	0		0	0		0
a			b			c	

rank(b, 5)

select(1, 2)



GMR


$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

1	0	0	1	0	0	1	0
a			b			c	

rank(b, 5)

select(1, 2)



GMR


$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

1	0	0	1	0	0	1	0
a			b			c	

$$\text{rank}(b, 5) = 2$$

select(1, 2)



GMR


$$S = [abcabbca]$$

$$\Pi = [1, 4, 8, 2, 5, 6, 3, 7]$$

1	0	0	1	0	0	1	0
a			b			c	

$$\text{rank}(b, 5) = 2$$

$\text{select}(1, 2)$



$O(\log \log \sigma)$ time

GMR



- It is possible to extend this to sequences of arbitrary length
- The resulting space is $n \log \sigma + n \cdot o(\log \sigma)$
- Times remain the same as for short sequences

Applications



k2-tree

- A space-efficient representation for graphs
- The main idea behind it is to reduce the graph to a tree
- We will discuss a simple version, assuming we represent the graph as a quad-tree

k2-tree




k2-tree



k2-tree


q_1				q_2
	1	0		
	1	1		
	1	0	1	1
q_4	0	0	1	1
				q_3




k2-tree

Diagram illustrating a k2-tree structure. The tree is represented as a 4x4 grid of cells, with rows labeled q_1 to q_4 and columns labeled q_2 to q_3 . The grid contains binary values (0 or 1) representing the tree structure. The grid is divided into four quadrants by a vertical line between columns 2 and 3, and a horizontal line between rows 2 and 3. The quadrants are labeled with large, faint numbers: 1 in the top-left, 0 in the top-right, 1 in the bottom-left, and 1 in the bottom-right.


q_1	q_2	q_3	q_4
1	0		
1	1		
1	0	1	1
0	0	1	1





k2-tree





k2-tree




k2-tree





k2-tree





k2-tree





k2-tree





k2-tree




k2-tree




k2-tree




k2-tree




k2-tree



k2-tree



k2-tree



k2trees

- Each node has either 4 or no children
- We can adapt the simple representation for binary trees we saw at the beginning
- This means, we need 1 bit per node!
- But how much is that compared to the information theoretical minimum?

k2tree


- There are $\binom{n^2}{m}$ graphs with m edges
- The number of bits we need to represent such an object is:

$$\log \binom{n^2}{m}$$

- This is roughly $\mathcal{H} = m \log \frac{n^2}{m} + O(m)$


k2tree

- Intuition: the worst that could happen
- All elements split their path as soon as possible
- We spend 4 bits per element per level after they split




k2tree

- Intuition: the worst that could happen
- All elements split their path as soon as possible
- We spend 4 bits per element per level after they split




k2tree

- Intuition: the worst that could happen
- All elements split their path as soon as possible
- We spend 4 bits per element per level after they split



k2tree


- Intuition: the worst that could happen
- All elements split their path as soon as possible
- We spend 4 bits per element per level after they split



k2tree

- Intuition: the worst that could happen
- All elements split their path as soon as possible
- We spend 4 bits per element per level after they split

$$4m \log_4 \frac{n^2}{m} + O(m)$$




k2tree

- Intuition: the worst that could happen
- All elements split their path as soon as possible
- We spend 4 bits per element per level after they split

$$4m \log_4 \frac{n^2}{m} + O(m)$$

$$2\mathcal{H} + O(m)$$




A Simple Trie

A simple trie

- A Trie is simply a tree symbols in the edges
- Very good option for representing string dictionaries


A simple trie

- A Trie is simply a tree symbols in the edges



A simple trie

- A Trie is simply a tree symbols in the edges




Finding an element: $O(m)$

A simple trie

- We will represent the tree using LOUDS
- Each l in the LOUDS representation will have a label associated with it
- We want to support the following queries:
 - $\text{child}(i, a)$: child of node i through label a , or -1 if it doesn't exist
 - $\text{parent-label}(i)$: which label brought me from my parent?

A simple trie

- A Trie is simply a tree symbols in the edges



A	E	I		A		A	E	A	E			A	E				
I	I	I	0	I	0	I	I	0	I	I	0	0	0	I	I	0	0
a				b		c		d		e	f	g		i	j	k	l

A simple trie

- Space required by the trie (in bits):

$$n \log \sigma + o(n \log \sigma)$$

- Time for label-related queries:

$$O(\log \sigma) \text{ or } O(\log \log \sigma)$$

A simple trie

- Space required by the trie (in bits):

$$n \log \sigma + o(n \log \sigma)$$

- Time for searching a string of length m

$$O(m \log \sigma)$$

- ... or

$$O(m \log \log \sigma)$$

Wrapping up

LIBCDS2

- And we are getting more help
 - Alex Bowe
 - Rodrigo Cánovas
 - Roberto Konow

LIBCDS2

- It's based on libcds, but for big datasets
- We are making some time tradeoffs, but it's easier to use
- It's almost ready to use in multi-threading settings
- Basic support for other languages (at the time Go, Python in process)

LIBCDS

- Version 1 -- “stable”
 - <http://libcds.recoded.cl>
- Version 2 -- ready to start trying it out
 - <http://libcds2.recoded.cl>
 - <http://github.com/fclaude/libcds2>

Conclusions

- We can save considerable space for static data structures using these techniques
- The same principles work in practice for dynamic data structures (but there is still a lot to be done here)
- Try it out and see for yourself how these structures run

References

- Bitmaps: Jacobson '89; Clark & Munro '96
- Huffman: Huffman '52
- LOUDS: Jacobson '89
- Wavelet Trees: Grossi, Gupta & Vitter '03
- Permutations: Munro, Raman, Raman & Rao '03
- GMR: Golynski, Munro & Rao '06
- K2Tree: Brisaboa, Ladra & Navarro '09
- DACs: Brisaboa, Ladra & Navarro '09

Thanks!



Now...

